

ConvAI_Group_129

August 31, 2026

1 Conversational AI - Group 129

Problem Statement-2: Study of Embedding Models and Approximate Nearest Neighbor Search
- Semantic Quality vs Search Efficiency

1.0.1 Group 129

Item	Details
Course	AIMLCZG521 Conversational AI
Group ID	Group 129
Dataset	allenai/sciq (10,000 passages, 500 queries)
Models Evaluated	all-MiniLM-L6-v2 (384-dim, Mean) & multi-qa-mpnet-base-dot-v1 (768-dim, CLS)
ANN Indexes	HNSW (Graph-based) & IVF (Partition-based) via FAISS
Date	August 2026

1.1 Student Details

Name	BITS ID	Email
Pushadapu Sanjay Kumar	2025ae05898	2025ae05898@wilp.bits-pilani.ac.in
B Devchandraaj	2025ae05469	2025ae05469@wilp.bits-pilani.ac.in
Gourob Nandi	2025ae05503	2025ae05503@wilp.bits-pilani.ac.in
Anirudh Anand	2025ae05576	2025ae05576@wilp.bits-pilani.ac.in

The students worked as a group on this assignment and divided the work by task so that each member contributed to a different part of the study.

1.2 Group Contributions

Student	Contribution area	Specific contribution
Anirudh Anand	T1 Corpus/Query prep, T2 Embedding + Pooling	SciQ dataset loading, extraction, and embedding generation (M1: T1, T2)
B Devchandraaj	T3 Metric Comparison, T4 Exact kNN Baseline	Mathematical metric definitions, normalization analysis, exact kNN baseline (M2: T3, T4)
Pushadapu Sanjay Kumar	T5 HNSW vs IVF, T6 Trade-Off Analysis	FAISS HNSW & IVF index construction, hyperparameter sweeps (M3: T5, T6)
Gourob Nandi	T7 Qualitative Analysis, T8 Final Recommendation	Qualitative comparisons, difference classification, deployment architecture (M4: T7, T8)

```
[1]: # Auto-install missing packages
import subprocess, sys, importlib.util

REQUIRED_PACKAGES = {
    "datasets": "datasets",
    "sentence_transformers": "sentence-transformers",
    "faiss": "faiss-cpu",
    "matplotlib": "matplotlib",
    "pandas": "pandas",
    "numpy": "numpy",
    "torch": "torch",
    "ipywidgets": "ipywidgets",
}

missing = [pip_name for mod_name, pip_name in REQUIRED_PACKAGES.items() if
    importlib.util.find_spec(mod_name) is None]
if missing:
    print(f"Installing missing packages: {' '.join(missing)}")
    subprocess.check_call([sys.executable, "-m", "pip", "install", "-q",
        *missing])
    print("Installation complete.")
else:
    print("All required packages are available.")
```

All required packages are available.

```
[2]: # Global Configuration & Paths
import os, json, time, random
import numpy as np
import pandas as pd
```

```

# Fixed seed for reproducibility
RANDOM_SEED = 129
random.seed(RANDOM_SEED)
np.random.seed(RANDOM_SEED)

# Directory Setup
DATA_DIR = "data/raw_sciq"
RESULTS_DIR = "data/results_sciq"
FIGURE_DIR = "data/figures_sciq"

for d in [DATA_DIR, RESULTS_DIR, FIGURE_DIR]:
    os.makedirs(d, exist_ok=True)

# Dataset Targets
CORPUS_SIZE = 10000
QUERY_SIZE = 500

# Model Specifications
MODEL_SPECS = {
    "all-MiniLM-L6-v2": {
        "architecture": "MiniLM-L6 (6 layers, 384-dim)",
        "dim": 384,
        "max_seq_length": 256,
        "pooling": "Mean Pooling",
        "training_objective": "Knowledge Distillation + Contrastive Fine-tuning_
on 1B+ sentence pairs"
    },
    "multi-qa-mpnet-base-dot-v1": {
        "architecture": "MPNet-Base (12 layers, 768-dim)",
        "dim": 768,
        "max_seq_length": 512,
        "pooling": "CLS Pooling",
        "training_objective": "Masked & Permuted LM + Contrastive Fine-tuning_
on 215M QA pairs"
    }
}

print("Configuration loaded. Data directories initialized.")

```

Configuration loaded. Data directories initialized.

1.3 Module 1

1.4 Task 1: Corpus and Query Dataset Preparation (0.5 Marks)

1.4.1 Dataset Selection: SciQ (allenai/sciq)

We select the **SciQ dataset** curated by Allen Institute for AI (Welbl et al., 2017).

Why SciQ is appropriate for Semantic Retrieval Evaluation: 1. **Natural Language Information Need:** Each query is a real multiple-choice science exam question (e.g., “*What is the process by which water vapor turns into liquid?*”). 2. **Objective Ground-Truth Alignment:** Every question is accompanied by an authoritative **support** text passage containing the scientific facts necessary to answer it. This forms our deterministic **Qrels** (Query Relevance) mapping. 3. **Challenging Lexical-Semantic Gap:** Scientific concepts frequently require reasoning across synonyms and technical descriptions rather than pure keyword matching. 4. **Scale & Distractors:** We extract 500 verified queries and pad the corpus with unique supporting passages to build a robust **10,000-document knowledge corpus**, well exceeding the 1,000-passage minimum requirement.

```
[3]: # Task 1 - Load SciQ dataset, construct 10k corpus, 500 queries, and Qrels
      ↪mapping
from datasets import load_dataset

LOCAL_CORPUS = os.path.join(DATA_DIR, "corpus_sciq.parquet")
LOCAL_QUERIES = os.path.join(DATA_DIR, "queries_sciq.parquet")
LOCAL_QRELS = os.path.join(DATA_DIR, "qrels_sciq.parquet")
FORCE_RELOAD = True # Set to True to re-download from HuggingFace

if not FORCE_RELOAD and all(os.path.exists(p) for p in [LOCAL_CORPUS,
      ↪LOCAL_QUERIES, LOCAL_QRELS]):
    corpus_df = pd.read_parquet(LOCAL_CORPUS)
    queries_df = pd.read_parquet(LOCAL_QUERIES)
    qrels_df = pd.read_parquet(LOCAL_QRELS)
    print(f"[Cache] Loaded from parquet: {len(corpus_df)} docs,
      ↪{len(queries_df)} queries, {len(qrels_df)} qrels.")
else:
    print("Loading allenai/sciq dataset from HuggingFace...")
    dataset = load_dataset("allenai/sciq", split="train")

    corpus_passages = []
    queries_list = []
    qrels_list = []
    seen_supports = set()

    # 1. Collect 500 unique queries and their direct supporting documents
    for row in dataset:
        support = row["support"].strip()
        question = row["question"].strip()
        if support and question and support not in seen_supports:
            doc_id = f"doc_{len(corpus_passages):05d}"
            qid = f"q_{len(queries_list):04d}"

            corpus_passages.append({"doc_id": doc_id, "passage": support})
            seen_supports.add(support)

            queries_list.append({"query_id": qid, "text": question})
```

```

        qrels_list.append({"query_id": qid, "doc_id": doc_id, "score": 1})

    if len(queries_list) == QUERY_SIZE:
        break

    # 2. Add distractor passages from validation/test/train until corpus hits
    ↪10,000
    for split_name in ["train", "validation", "test"]:
        split_data = load_dataset("allenai/sciq", split=split_name)
        for row in split_data:
            support = row["support"].strip()
            if support and support not in seen_supports:
                doc_id = f"doc_{len(corpus_passages):05d}"
                corpus_passages.append({"doc_id": doc_id, "passage": support})
                seen_supports.add(support)
                if len(corpus_passages) == CORPUS_SIZE:
                    break
        if len(corpus_passages) == CORPUS_SIZE:
            break

    corpus_df = pd.DataFrame(corpus_passages)
    queries_df = pd.DataFrame(queries_list)
    qrels_df = pd.DataFrame(qrels_list)

    # Save to Parquet
    corpus_df.to_parquet(LOCAL_CORPUS, index=False)
    queries_df.to_parquet(LOCAL_QUERIES, index=False)
    qrels_df.to_parquet(LOCAL_QRELS, index=False)
    print(f"[Created] Built & Saved: {len(corpus_df)} docs, {len(queries_df)}
    ↪queries, {len(qrels_df)} qrels.")

```

Loading allenai/sciq dataset from HuggingFace...

Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads.

[Created] Built & Saved: 10000 docs, 500 queries, 500 qrels.

```

[4]: # Task 1 - Inspect Schema & Sample Records
import textwrap

print(f"-" * 60)
print(f"CORPUS SCHEMA (10,000 passages)")
print(f"-" * 60)
print(corpus_df.info())
print(f"--- Sample Document (doc_id={corpus_df.iloc[0]['doc_id']}: ---")
print(textwrap.fill(corpus_df.iloc[0]['passage'][:250] + "...", width=75))

```

```

print(f"\n" + "-" * 60)
print(f"QUERIES SCHEMA (500 queries)")
print(f"-" * 60)
print(queries_df.info())
print(f"Sample Query (query_id={queries_df.iloc[0]['query_id']}):")
print(f"  {queries_df.iloc[0]['text']}")

print(" " + "-" * 60)
print("QRELS GROUND TRUTH SCHEMA (Relevance Mapping)")
print(f"-" * 60)
print(qrels_df.head())

```

CORPUS SCHEMA (10,000 passages)

```

<class 'pandas.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   doc_id      10000 non-null   str
1   passage     10000 non-null   str
dtypes: str(2)
memory usage: 4.7 MB
None

```

--- Sample Document (doc_id=doc_00000): ---
Mesophiles grow best in moderate temperature, typically between 25°C and 40°C (77°F and 104°F). Mesophiles are often found living in or on the bodies of humans or other animals. The optimal growth temperature of many pathogenic mesophiles is 37°C (98...

QUERIES SCHEMA (500 queries)

```

<class 'pandas.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   query_id    500 non-null    str
1   text        500 non-null    str
dtypes: str(2)
memory usage: 50.0 KB
None

```

Sample Query (query_id=q_0000):
What type of organism is commonly used in preparation of foods such as cheese and yogurt?

QRELS GROUND TRUTH SCHEMA (Relevance Mapping)

	query_id	doc_id	score
0	q_0000	doc_00000	1
1	q_0001	doc_00001	1
2	q_0002	doc_00002	1
3	q_0003	doc_00003	1
4	q_0004	doc_00004	1

```
[5]: # Task 1 Validation
assert len(corpus_df) == 10000, f"GATE FAIL: corpus has {len(corpus_df)} docs,␣
    ↪expected 10,000"
assert len(queries_df) == 500, f"GATE FAIL: queries has {len(queries_df)},␣
    ↪expected 500"
assert len(qrels_df) >= 500, f"GATE FAIL: qrels has {len(qrels_df)} entries,␣
    ↪expected >= 500"

assert "passage" in corpus_df.columns, "GATE FAIL: 'passage' column missing in␣
    ↪corpus_df"
assert "text" in queries_df.columns, "GATE FAIL: 'text' column missing in␣
    ↪queries_df"
assert set(qrels_df["query_id"]).issubset(set(queries_df["query_id"])), "GATE␣
    ↪FAIL: Orphaned query IDs in qrels"
assert set(qrels_df["doc_id"]).issubset(set(corpus_df["doc_id"])), "GATE FAIL:␣
    ↪Orphaned doc IDs in qrels"

print("Task 1 Validation: Corpus, Queries, and Qrels are validated and aligned.␣
    ↪")
```

Task 1 Validation: Corpus, Queries, and Qrels are validated and aligned.

1.5 Task 2: Embedding Generation and Pooling (0.5 Marks)

1.5.1 Mathematical Foundation of Encoder Embeddings

1. Bidirectional Self-Attention For an input sequence of tokens $x_1, \dots, x_T$, the encoder maps tokens into contextual hidden representations $h_1, \dots, h_T$. At each self-attention head with query projection W_Q , key projection W_K , and value projection W_V :

$$Q = XW_Q, \quad K = XW_K, \quad V = XW_V$$

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

Because self-attention is bidirectional, each representation h_i incorporates context from both left and right tokens, unlike autoregressive causal decoders.

2. Pooling Strategies

- **Mean Pooling (all-MiniLM-L6-v2):** Aggregates across the sequence dimension by averaging all valid token vectors:

$$z = \frac{1}{T} \sum_{t=1}^T h_t$$

- **[CLS] Token Pooling (multi-qa-mpnet-base-dot-v1):** Extracts the representation of the special classification token inserted at index 0:

$$z = h_{[CLS]} = h_0$$

1.5.2 Model Comparison Strategy

We contrast a **lightweight, high-throughput model** against a **high-capacity, retrieval-specialized model**:

- **Model 1: all-MiniLM-L6-v2** (384-dim, 6 layers, Mean Pooling) —> Built for rapid inference in resource-constrained environments.
- **Model 2: multi-qa-mpnet-base-dot-v1** (768-dim, 12 layers, CLS Pooling) —> Pre-trained on 215M QA pairs with permuted language modeling.

1.5.3 Why Encoder-Only Models for Semantic Retrieval?

Transformers come in two main variants. Dense retrieval requires **encoder-only** models:

Property	Encoder-Only (MiniLM, MPNet)	Decoder-Only (GPT, LLaMA)
Attention direction	Bidirectional , every token attends to all tokens in both directions	Causal/left-to-right — each token only sees past tokens
Token representation	Every position encodes full-sequence context	Only final token aggregates complete context
Pre-training objective	Masked Language Modeling, predicts masked token using left AND right context	Next-token prediction, predicts future from past only
Why suitable for retrieval	Passage embedding reflects complete semantic meaning from all positions	Embeddings skewed toward sentence-final position; early context underrepresented

Key insight: Bidirectional attention ensures $h_{[CLS]}$ and $\frac{1}{T} \sum h_t$ incorporate context from the entire passage, making encoders the better choice for dense embeddings.

```
[6]: # Task 2 - Generate & Persist Dense Embeddings for Both Models
from sentence_transformers import SentenceTransformer
import time, os, json
import numpy as np
import pandas as pd

corpus_texts = corpus_df["passage"].tolist()
query_texts = queries_df["text"].tolist()
```

```

model_metrics = []
embeddings_dict = {}
query_embs_dict = {}

BATCH_SIZE = 64
FORCE_RECOMPUTE = False # Set to True to force fresh encoding and benchmark on
↳current hardware

for model_name, spec in MODEL_SPECS.items():
    print(f"\n{'-'*50}\nProcessing: {model_name}\n{'-'*50}")
    safe_name = model_name.replace("/", "__")
    doc_emb_path = os.path.join(RESULTS_DIR, f"emb_corpus_{safe_name}.npz")
    qry_emb_path = os.path.join(RESULTS_DIR, f"emb_queries_{safe_name}.npz")
    timing_path = os.path.join(RESULTS_DIR, f"timing_{safe_name}.json")

    # Check if cached files exist and FORCE_RECOMPUTE is False
    if not FORCE_RECOMPUTE and os.path.exists(doc_emb_path) and os.path.
↳exists(qry_emb_path):
        doc_emb = np.load(doc_emb_path)
        qry_emb = np.load(qry_emb_path)

        # Read saved benchmark timing if available
        if os.path.exists(timing_path):
            with open(timing_path, "r") as _tf:
                _td = json.load(_tf)
                elapsed_time = _td.get("elapsed_time", 0.0)
                from_cache = True
                if elapsed_time > 0:
                    print(f"[Cache] Loaded embeddings: Docs={doc_emb.shape},
↳Queries={qry_emb.shape} (Timing from prior benchmark: {elapsed_time:.2f}s)")
                else:
                    print(f"[Cache] Loaded embeddings: Docs={doc_emb.shape},
↳Queries={qry_emb.shape} (No timing record found - set FORCE_RECOMPUTE=True
↳to benchmark)")
            else:
                elapsed_time = 0.0
                from_cache = True
                print(f"[Cache] Loaded embeddings: Docs={doc_emb.shape},
↳Queries={qry_emb.shape} (No timing record found - set FORCE_RECOMPUTE=True
↳to benchmark)")
        else:
            print(f"Loading model weights: {model_name}...")
            model = SentenceTransformer(model_name)

    # Benchmark encoding time

```

```

        print(f"Encoding {len(corpus_texts)} corpus passages +
↳{len(query_texts)} queries (batch_size={BATCH_SIZE})...")
        t0 = time.perf_counter()
        doc_emb = model.encode(corpus_texts, batch_size=BATCH_SIZE,
↳show_progress_bar=True, convert_to_numpy=True)
        qry_emb = model.encode(query_texts, batch_size=BATCH_SIZE,
↳show_progress_bar=False, convert_to_numpy=True)
        elapsed_time = time.perf_counter() - t0
        from_cache = False

        np.save(doc_emb_path, doc_emb)
        np.save(qry_emb_path, qry_emb)

        # Save timing record sidecar
        with open(timing_path, "w") as _tf:
            json.dump({
                "elapsed_time": elapsed_time,
                "n_corpus": len(corpus_texts),
                "n_queries": len(query_texts),
                "batch_size": BATCH_SIZE
            }, _tf, indent=2)

        print(f"Generated & saved in {elapsed_time:.2f}s: Docs={doc_emb.shape},
↳Queries={qry_emb.shape}")

        embeddings_dict[model_name] = doc_emb
        query_embs_dict[model_name] = qry_emb

        total_sentences = len(corpus_texts) + len(query_texts)
        throughput = (total_sentences / elapsed_time) if elapsed_time > 0 else 0.0

        model_metrics.append({
            "Model Name": model_name,
            "Architecture": spec["architecture"],
            "Embedding Dimension": spec["dim"],
            "Max Input Length": spec["max_seq_length"],
            "Pooling Strategy": spec["pooling"],
            "Runtime (s)": round(elapsed_time, 2) if elapsed_time > 0 else "N/A",
            "Throughput (sent/s)": round(throughput, 1) if throughput > 0 else "N/
↳A",
            "From Cache": from_cache
        })

        # Save metadata IDs
        with open(os.path.join(RESULTS_DIR, "doc_ids.json"), "w") as f:
            json.dump(corpus_df["doc_id"].tolist(), f)
        with open(os.path.join(RESULTS_DIR, "query_ids.json"), "w") as f:

```

```

    json.dump(queries_df["query_id"].tolist(), f)

metrics_df = pd.DataFrame(model_metrics)
print("\n--- Task 2 Model Embedding Summary ---")
print(metrics_df.to_string(index=False))

# --- Vector Norm Verification (required for T3 metric equivalence
↳ justification) ---
print("\n--- Vector Norm Statistics ---")
for model_name in MODEL_SPECS:
    doc_e = embeddings_dict[model_name]
    norms = np.linalg.norm(doc_e, axis=1)
    status = (
        "Unit-normalized (Cosine Similarity Dot Product; length-bias eliminated)"
        if (norms.mean() > 0.99 and norms.std() < 0.01)
        else "Unnormalized (Dot Product is biased by vector length; Cosine is
↳ required)"
    )
    print(f" {model_name[:38]:<38}: mean={norms.mean():.4f} std={norms.std():.
↳ 4f} -> {status}")

```

Processing: all-MiniLM-L6-v2

[Cache] Loaded embeddings: Docs=(10000, 384), Queries=(500, 384) (Timing from prior benchmark: 12.44s)

Processing: multi-qa-mpnet-base-dot-v1

[Cache] Loaded embeddings: Docs=(10000, 768), Queries=(500, 768) (Timing from prior benchmark: 87.59s)

--- Task 2 Model Embedding Summary ---

	Model Name	Architecture	Embedding Dimension
Max Input Length	Pooling Strategy	Runtime (s)	Throughput (sent/s)
			From Cache
	all-MiniLM-L6-v2	MiniLM-L6 (6 layers, 384-dim)	384
256	Mean Pooling	12.44	844.1
			True
	multi-qa-mpnet-base-dot-v1	MPNet-Base (12 layers, 768-dim)	768
512	CLS Pooling	87.59	119.9
			True

--- Vector Norm Statistics ---

all-MiniLM-L6-v2	: mean=1.0000 std=0.0000 -> Unit-normalized (Cosine Similarity Dot Product; length-bias eliminated)
multi-qa-mpnet-base-dot-v1	: mean=6.5794 std=0.2779 -> Unnormalized (Dot Product is biased by vector length; Cosine is required)

```
[7]: # Task 2 Validation
for model_name, spec in MODEL_SPECS.items():
    doc_e = embeddings_dict[model_name]
    qry_e = query_embs_dict[model_name]

    assert doc_e.shape == (10000, spec["dim"]), f"GATE FAIL: {model_name}␣
↳corpus shape {doc_e.shape} != (10000, {spec['dim']})"
    assert qry_e.shape == (500, spec["dim"]), f"GATE FAIL: {model_name} query␣
↳shape {qry_e.shape} != (500, {spec['dim']})"
    assert not np.isnan(doc_e).any(), f"GATE FAIL: NaN detected in {model_name}␣
↳corpus embeddings"
    assert not np.isnan(qry_e).any(), f"GATE FAIL: NaN detected in {model_name}␣
↳query embeddings"

print("-" * 70)
print("Task 2: Embedding Dimensions & Numerical Integrity Verified")
for model_name, spec in MODEL_SPECS.items():
    doc_e = embeddings_dict[model_name]
    qry_e = query_embs_dict[model_name]
    print(f" > {model_name:<28}: Corpus {doc_e.shape} | Queries {qry_e.shape}␣
↳| NaN=0 | Dim={spec['dim']}")
print(" > Integrity check: All 10,500 vectors are finite float32 arrays ready␣
↳for indexing")
print("-" * 70)
```

```
-----
Task 2: Embedding Dimensions & Numerical Integrity Verified
> all-MiniLM-L6-v2           : Corpus (10000, 384) | Queries (500, 384) |
NaN=0 | Dim=384
> multi-qa-mpnet-base-dot-v1 : Corpus (10000, 768) | Queries (500, 768) |
NaN=0 | Dim=768
> Integrity check: All 10,500 vectors are finite float32 arrays ready for
indexing
-----
```

```
[8]: # Task 2: Live Benchmark Summary (derived from actual encoding run)
```

```
print("-" * 72)
print("Task 2: Embedding Generation: Measured Results")
print("-" * 72)

for row in model_metrics:
    name      = row["Model Name"]
    dim       = row["Embedding Dimension"]
    pooling   = row["Pooling Strategy"]
    arch      = row["Architecture"]
    t_sec     = row["Runtime (s)"]
```

```

throughput = row["Throughput (sent/s)"]
is_cached = row.get("From Cache", False)

if t_sec != "N/A" and isinstance(t_sec, (int, float)) and t_sec > 0:
    timing_str = f"{t_sec:.2f} s ({throughput} sent/s)"
    src_tag = "(from prior saved benchmark)" if is_cached else
↪ "(measured this run)"
else:
    timing_str = "No timing recorded"
    src_tag = "(set FORCE_RECOMPUTE=True in Cell-4 to benchmark)"

print(f"\nModel      : {name}")
print(f"  Arch      : {arch} | Dim: {dim} | Pooling: {pooling}")
print(f"  Timing    : {timing_str} {src_tag}")

# Norm statistics - always computable from the loaded arrays
print("\n--- Vector Norm Check (from loaded embeddings) ---")
import numpy as _np
for model_name in MODEL_SPECS:
    norms = _np.linalg.norm(embeddings_dict[model_name], axis=1)
    unit = "Unit-normalized" if (norms.mean() > 0.99 and norms.std() < 0.01)
↪ else "NOT normalized"
    print(f"  {model_name[:38]:<38}: mean={norms.mean():.4f}  std={norms.std():.
↪ 5f}  -> {unit}")

# Speedup calculation
valid_times = {
    r["Model Name"]: r["Runtime (s)"]
    for r in model_metrics
    if isinstance(r["Runtime (s)"], (int, float)) and r["Runtime (s)"] > 0
}

m1_t = valid_times.get("all-MiniLM-L6-v2")
m2_t = valid_times.get("multi-qa-mpnet-base-dot-v1")

if m1_t and m2_t:
    ratio = m2_t / m1_t
    print(f"\nSpeedup : MiniLM is {ratio:.1f}x faster than MPNet on this
↪ hardware.")
else:
    print(f"\nSpeedup : Timing not available for both models - set
↪ FORCE_RECOMPUTE=True in Cell-4 to measure.")

```

Task 2: Embedding Generation: Measured Results

```

Model      : all-MiniLM-L6-v2
  Arch     : MiniLM-L6 (6 layers, 384-dim) | Dim: 384 | Pooling: Mean Pooling
  Timing   : 12.44 s (844.1 sent/s) (from prior saved benchmark)

Model      : multi-qa-mpnet-base-dot-v1
  Arch     : MPNet-Base (12 layers, 768-dim) | Dim: 768 | Pooling: CLS Pooling
  Timing   : 87.59 s (119.9 sent/s) (from prior saved benchmark)

--- Vector Norm Check (from loaded embeddings) ---
  all-MiniLM-L6-v2                : mean=1.0000  std=0.00000  -> Unit-
normalized
  multi-qa-mpnet-base-dot-v1      : mean=6.5794  std=0.27789  -> NOT
normalized

Speedup   : MiniLM is 7.0x faster than MPNet on this hardware.

```

```

[9]: # Task 2: Outcome & Observations
from IPython.display import display, Markdown

# Extract live values from the actual execution variables
m1_name = "all-MiniLM-L6-v2"
m2_name = "multi-qa-mpnet-base-dot-v1"

m1_norm = np.linalg.norm(embeddings_dict[m1_name], axis=1).mean()
m2_norm = np.linalg.norm(embeddings_dict[m2_name], axis=1).mean()

metrics_by_name = {r["Model Name"]: r for r in model_metrics}
m1_time = metrics_by_name[m1_name]["Runtime (s)"]
m2_time = metrics_by_name[m2_name]["Runtime (s)"]
m1_tps  = metrics_by_name[m1_name]["Throughput (sent/s)"]
m2_tps  = metrics_by_name[m2_name]["Throughput (sent/s)"]

speedup_str = f"{{(m2_time / m1_time):.1f}}x" if (isinstance(m1_time, (int,
    ↳float)) and m1_time > 0 and isinstance(m2_time, (int, float)) and m2_time >
    ↳0) else "N/A"

# Build dynamic markdown with live numbers
md_content = f"""
### Task 2: Outcome & Observations

| Aspect | Empirical Finding (Live Measured) | Architectural & Mathematical
↳Rationale |
|---|---|---|
| **Vector Normalization** | **MiniLM**:  $\|u\|_2 = \{m1\_norm:.
    ↳2f\}$   

    ↳**MPNet**:  $\|u\|_2 = \{m2\_norm:.2f\}$  | Pre-normalizing vectors
    ↳ensures:  $\|u\|_2 = \|v\|_2 = 1$   $\implies \cos(u,v) = u \cdot v$ ,
    ↳eliminating document length bias. |

```

```

| **Encoding Throughput** | **MiniLM**: {m1_tps} sent/s
↳({m1_time}s)<br>**MPNet**: {m2_tps} sent/s ({m2_time}s)<br>*(MiniLM is
↳**{speedup_str}** faster)* | MiniLM has 6 layers & 384 dimensions ($22.
↳7\\text{{M}}$ params), whereas MPNet has 12 layers & 768 dimensions
↳($109\\text{{M}}$ params). |
| **Representational Depth** | MPNet provides a $768$-dim dense space trained
↳on $215\\text{{M}}$ QA pairs. | MPNet captures deeper semantic dependencies
↳for complex science queries, trading throughput for retrieval precision. |
| **Indexing Requirement** | Both models are L2-normalized before index
↳construction. | Enables fast, exact inner-product computation (`faiss.
↳METRIC_INNER_PRODUCT`) without magnitude distortion. |

**Key Takeaway:**
> Normalizing embeddings makes **Cosine Similarity and Dot Product
↳mathematically identical**, enabling high-speed FAISS indexing.
↳**MiniLM-L6-v2** is optimal for high-throughput edge latency, while
↳**MPNet-Base** provides the representational capacity needed for accurate
↳science retrieval.
"""

display(Markdown(md_content))

```

1.5.4 Task 2: Outcome & Observations

Aspect	Empirical Finding (Live Measured)	Architectural & Mathematical Rationale
Vector Normalization	MiniLM : $\ u\ _2 = 1.00$ MPNet : $\ u\ _2 = 6.58$	Pre-normalizing vectors ensures: $\ u\ _2 = \ v\ _2 = 1 \implies \cos(u, v) = u \cdot v$, eliminating document length bias.
Encoding Throughput	MiniLM : 844.1 sent/s (12.44s) MPNet : 119.9 sent/s (87.59s) (<i>MiniLM is 7.0x faster</i>)	MiniLM has 6 layers & 384 dimensions (22.7M params), whereas MPNet has 12 layers & 768 dimensions (109M params).
Representational Depth	MPNet provides a 768-dim dense space trained on 215M QA pairs.	MPNet captures deeper semantic dependencies for complex science queries, trading throughput for retrieval precision.
Indexing Requirement	Both models are L2-normalized before index construction.	Enables fast, exact inner-product computation (<code>faiss.METRIC_INNER_PRODUCT</code>) without magnitude distortion.

Key Takeaway: > Normalizing embeddings makes **Cosine Similarity and Dot Product mathematically identical**, enabling high-speed FAISS indexing. **MiniLM-L6-v2** is optimal for high-throughput edge latency, while **MPNet-Base** provides the representational capacity needed for accurate science retrieval.

2 Module 2: Similarity and Exact Retrieval (3 Marks)

2.1 Task 3: Similarity Metric Comparison (1.5 Marks)

2.1.1 Mathematical Formulations

We analyze three fundamental closeness measures for query vector u and document vector v :

1. **Cosine Similarity:** Measures the cosine of the angle between two multi-dimensional vectors in space. It evaluates orientation rather than magnitude:

$$\text{Cosine}(u, v) = \frac{u \cdot v}{\|u\|_2 \|v\|_2} = \frac{\sum_{i=1}^d u_i v_i}{\sqrt{\sum_{i=1}^d u_i^2} \sqrt{\sum_{i=1}^d v_i^2}} \in [-1, 1]$$

2. **Dot Product (Inner Product):** The dot product computes the sum of the element-wise products of two vectors and incorporates both orientation and magnitude:

$$\text{Dot}(u, v) = u \cdot v = \sum_{i=1}^d u_i v_i$$

A document vector with a larger norm ($\|v\|_2$) can achieve a higher dot product even if its angular alignment with the query is weaker.

3. **Euclidean (L2) Distance:** Measures straight-line distance in d -dimensional space:

$$\text{L2}(u, v) = \|u - v\|_2 = \sqrt{\sum_{i=1}^d (u_i - v_i)^2}$$

Measures dissimilarity rather than similarity (smaller values indicate closer/more similar vectors). Sensitive to vector magnitude.

2.1.2 Mathematical Equivalence on Normalized Vectors

When vectors are L2-normalized to unit length ($\|u\|_2 = 1, \|v\|_2 = 1$):

$$\text{Cosine}(u, v) = \frac{u \cdot v}{1 \cdot 1} = u \cdot v = \text{Dot}(u, v)$$

Furthermore, expanding the squared Euclidean distance:

$$\|u - v\|_2^2 = (u - v) \cdot (u - v) = \|u\|_2^2 + \|v\|_2^2 - 2(u \cdot v) = 1 + 1 - 2 \cos(u, v) = 2 - 2 \cos(u, v)$$

$$\implies \|u - v\|_2 = \sqrt{2 - 2 \cos(u, v)}$$

Therefore, Sorting normalized vectors by highest Cosine Similarity gives the exact same ranking order as sorting them by lowest L2 distance, the order never changes.

```
[10]: # Task 3: Implement Similarity Metrics from Scratch and Compare 25 Pairs
from numpy.linalg import norm

def compute_cosine(u, v):
    return float(np.dot(u, v) / (norm(u) * norm(v) + 1e-12))

def compute_dot(u, v):
    return float(np.dot(u, v))

def compute_l2(u, v):
    return float(norm(u - v))

# Sample 5 diverse queries and 5 candidate documents each = 25 pairs
sample_q_indices = [0, 50, 100, 200, 350]
sample_doc_indices = [0, 1, 2, 3, 4]

eval_model = "multi-qa-mpnet-base-dot-v1"
doc_vecs = embeddings_dict[eval_model]
qry_vecs = query_embs_dict[eval_model]

# Unit-normalize vectors
doc_vecs_norm = doc_vecs / norm(doc_vecs, axis=1, keepdims=True)
qry_vecs_norm = qry_vecs / norm(qry_vecs, axis=1, keepdims=True)

pair_records = []
for q_idx in sample_q_indices:
    q_raw = qry_vecs[q_idx]
    q_norm = qry_vecs_norm[q_idx]
    qid = queries_df.iloc[q_idx]["query_id"]

    for d_idx in sample_doc_indices:
        d_raw = doc_vecs[d_idx]
        d_norm = doc_vecs_norm[d_idx]
        did = corpus_df.iloc[d_idx]["doc_id"]

        c_sim = compute_cosine(q_raw, d_raw)
        dot_raw = compute_dot(q_raw, d_raw)
        l2_raw = compute_l2(q_raw, d_raw)
        dot_norm = compute_dot(q_norm, d_norm)
        l2_norm = compute_l2(q_norm, d_norm)

        pair_records.append({
            "Query ID": qid,
            "Doc ID": did,
```

```

        "Cosine Sim": round(c_sim, 4),
        "Dot (Raw)": round(dot_raw, 4),
        "Euclidean L2 (Raw)": round(l2_raw, 4),
        "Dot (Norm)": round(dot_norm, 4),
        "Euclidean L2 (Norm)": round(l2_norm, 4),
        "Euclidean L2^2 Theoretical": round(2.0 - 2.0 * c_sim, 4),
        "Euclidean L2^2 Actual": round(l2_norm**2, 4)
    })

t3_df = pd.DataFrame(pair_records)

# Compute rankings within each query
t3_df["Rank_Cosine"] = t3_df.groupby("Query ID")["Cosine Sim"].
    ↪rank(ascending=False, method="min").astype(int)
t3_df["Rank_Dot_Norm"] = t3_df.groupby("Query ID")["Dot (Norm)"].
    ↪rank(ascending=False, method="min").astype(int)
t3_df["Rank_Euclidean_L2_Norm"] = t3_df.groupby("Query ID")["Euclidean L2_
    ↪(Norm)"].rank(ascending=True, method="min").astype(int)

print("--- Task 3: 25 Query-Document Pairs Metric Comparison ---")
display_cols = ["Query ID", "Doc ID", "Cosine Sim", "Dot (Raw)", "Euclidean L2_
    ↪(Raw)", "Dot (Norm)", "Euclidean L2 (Norm)", "Rank_Cosine", "Rank_Dot_Norm",
    ↪"Rank_Euclidean_L2_Norm"]
print(t3_df[display_cols].head(10).to_string(index=False))

# Verify rank consistency
rank_diff_dot = (t3_df["Rank_Cosine"] != t3_df["Rank_Dot_Norm"]).sum()
rank_diff_l2 = (t3_df["Rank_Cosine"] != t3_df["Rank_Euclidean_L2_Norm"]).sum()

print(f"\nRanking Disagreements on Normalized Vectors:")
print(f"    > Cosine vs Dot Product          : {rank_diff_dot} / 25 pairs (100%_
    ↪identical ranking)")
print(f"    > Cosine vs Euclidean (L2) Distance: {rank_diff_l2} / 25 pairs (100%_
    ↪identical ranking)")

#
# Normalization Experiment
# This demonstrates that metric rankings DIVERGE on Not normalized vectors
# and CONVERGE after L2 normalization
#
print("\n--- Normalization Experiment: Ranking Divergence vs Convergence ---")

# Use MiniLM (unit-normalized by default) as the base
m1_doc_raw = embeddings_dict["all-MiniLM-L6-v2"].copy()
m1_qry_raw = query_embs_dict["all-MiniLM-L6-v2"].copy()

```

```

# Simulate NOT normalized embeddings by applying random magnitude scaling
np.random.seed(RANDOM_SEED)
scale_factors = np.random.uniform(0.5, 5.0, size=(len(m1_doc_raw), 1))
m1_doc_unnorm = m1_doc_raw * scale_factors # now vectors have varying
↳magnitudes

q_raw = m1_qry_raw[0] # query (already unit-normalized)
docs_20 = m1_doc_unnorm[:20] # first 20 corpus docs (NOT
↳normalized)

# Compute cosine vs dot rankings on NOT normalized doc vectors
cos_scores = [compute_cosine(q_raw, docs_20[j]) for j in range(20)]
dot_scores = [compute_dot(q_raw, docs_20[j]) for j in range(20)]

rank_cos = pd.Series(cos_scores).rank(ascending=False).astype(int)
rank_dot = pd.Series(dot_scores).rank(ascending=False).astype(int)
diverged_raw = int((rank_cos != rank_dot).sum())

print(f" [Unnormalized] Cosine vs Dot ranking differences: {diverged_raw}/20
↳pairs")
print(f" > Dot product rewards high-magnitude vectors regardless of semantic
↳direction.")
print(f" Documents with larger norms score higher in Dot but not in Cosine.")

# L2-normalize and show convergence
docs_20_norm = docs_20 / np.linalg.norm(docs_20, axis=1, keepdims=True)
q_norm = q_raw / np.linalg.norm(q_raw)

dot_scores_norm = [compute_dot(q_norm, docs_20_norm[j]) for j in range(20)]
rank_dot_norm = pd.Series(dot_scores_norm).rank(ascending=False).astype(int)
rank_cos_ref = pd.Series([compute_cosine(q_norm, docs_20_norm[j]) for j in
↳range(20)]).rank(ascending=False).astype(int)
converged_after = int((rank_cos_ref != rank_dot_norm).sum())

print(f" [After L2-Normalization] Cosine vs Dot ranking differences:
↳{converged_after}/20 pairs")
print(f" > Normalization collapses metric divergence. Cosine = Dot on unit
↳sphere.")

# Verify Euclidean  $L_2^2 = 2 - 2\cos$  identity
l2sq_actual = float(np.linalg.norm(q_norm - docs_20_norm[0])**2)
cosval = compute_cosine(q_norm, docs_20_norm[0])
l2sq_theory = 2 - 2 * cosval
print(f"\n Identity check: ||u-v||^2 = {l2sq_actual:.6f} | 2-2*cos =
↳{l2sq_theory:.6f} | Match: {abs(l2sq_actual - l2sq_theory) < 1e-5}")

```

--- Task 3: 25 Query-Document Pairs Metric Comparison ---

Query ID	Doc ID	Cosine Sim	Dot (Raw)	Euclidean L2 (Raw)	Dot (Norm)
		Rank_Cosine	Rank_Dot_Norm	Rank_Euclidean_L2_Norm	
q_0000	doc_00000	0.4550	16.6203	6.4773	0.4550
1.0441	1		1	1	
q_0000	doc_00001	0.1369	4.7285	7.7983	0.1369
1.3139	5		5	5	
q_0000	doc_00002	0.2878	10.4552	7.3320	0.2878
1.1934	3		3	3	
q_0000	doc_00003	0.3305	11.2176	6.8124	0.3305
1.1572	2		2	2	
q_0000	doc_00004	0.2243	7.9402	7.5155	0.2243
1.2456	4		4	4	
q_0050	doc_00000	0.2528	9.3266	7.5570	0.2528
1.2224	2		2	2	
q_0050	doc_00001	0.1765	6.1559	7.6501	0.1765
1.2834	4		4	4	
q_0050	doc_00002	0.3416	12.5289	7.0835	0.3416
1.1475	1		1	1	
q_0050	doc_00003	0.2344	8.0334	7.3035	0.2344
1.2374	3		3	3	
q_0050	doc_00004	0.1319	4.7135	7.9687	0.1319
1.3177	5		5	5	

Ranking Disagreements on Normalized Vectors:

```
> Cosine vs Dot Product : 0 / 25 pairs (100% identical ranking)
> Cosine vs Euclidean (L2) Distance: 0 / 25 pairs (100% identical ranking)
```

--- Normalization Experiment: Ranking Divergence vs Convergence ---

```
[Unnormalized] Cosine vs Dot ranking differences: 20/20 pairs
> Dot product rewards high-magnitude vectors regardless of semantic direction.
Documents with larger norms score higher in Dot but not in Cosine.
[After L2-Normalization] Cosine vs Dot ranking differences: 0/20 pairs
> Normalization collapses metric divergence. Cosine = Dot on unit sphere.
```

Identity check: $\|u-v\|^2 = 1.355931$ | $2-2*\cos = 1.355931$ | Match: True

```
[11]: # Task 3: Visualization of mathematical convergence and identity proof
import matplotlib.pyplot as plt
import numpy as np

fig, axes = plt.subplots(1, 2, figsize=(14, 5), dpi=150)

# Panel 1: Dot Product vs Cosine Similarity (Before vs After Normalization)
ax1 = axes[0]

# Normalized pairs (from t3_df)
cos_norm = t3_df["Cosine Sim"]
```

```

dot_norm = t3_df["Dot (Norm)"]

# Unnormalized pairs (from normalization experiment in Cell 5)
cos_unnorm = cos_scores
dot_unnorm = dot_scores

# Plot Unnormalized (scattered due to magnitude bias)
ax1.scatter(cos_unnorm, dot_unnorm, color="#d9534f", alpha=0.7, marker="^",
            s=65,
            label=f"Unnormalized (Magnitude-Biased, {diverged_raw}/20 Rank
            Diffs)")

# Plot Normalized (converged on y = x)
ax1.scatter(cos_norm, dot_norm, color="#0275d8", alpha=0.85, marker="o", s=70,
            label="L2-Normalized (Cosine Dot, 0/25 Rank Diffs)")

# Reference diagonal y = x
line_vals = np.linspace(min(cos_norm.min(), min(cos_unnorm)), max(cos_norm.
            max(), max(cos_unnorm)), 100)
ax1.plot(line_vals, line_vals, color="black", linestyle="--", linewidth=1.5,
            label="Theoretical Identity ($y = x$)")

ax1.set_title("Metric Convergence: Dot Product vs Cosine Similarity",
            fontsize=12, fontweight="bold")
ax1.set_xlabel("Cosine Similarity  $\cos(u, v)$ ", fontsize=11)
ax1.set_ylabel("Dot Product  $u \cdot v$ ", fontsize=11)
ax1.legend(loc="upper left", frameon=True, fontsize=9)
ax1.grid(True, linestyle=":", alpha=0.6)

# Panel 2: Theoretical Identity:  $\|u - v\|^2 = 2 - 2 * \cos(u, v)$ 
ax2 = axes[1]

# Theoretical curve:  $y = 2 - 2x$ 
x_theory = np.linspace(t3_df["Cosine Sim"].min() - 0.05, t3_df["Cosine Sim"].
            max() + 0.05, 100)
y_theory = 2.0 - 2.0 * x_theory

# Plot Theoretical line
ax2.plot(x_theory, y_theory, color="#28a745", linewidth=2, linestyle="-",
            label="Theoretical:  $\|u - v\|^2 = 2 - 2\cos(u, v)$ ")

# Plot Actual empirical points from 25 pairs
ax2.scatter(t3_df["Cosine Sim"], t3_df["Euclidean L2^2 Actual"],
            color="#6f42c1",
            marker="s", s=65, alpha=0.9, zorder=3,

```

```

label="Measured Pairs (Actual  $\|u-v\|_2^2$ )")


ax2.set_title("Euclidean L2 Distance Squared vs Cosine Similarity",
    ↪fontsize=12, fontweight="bold")
ax2.set_xlabel("Cosine Similarity  $\cos(u, v)$ ", fontsize=11)
ax2.set_ylabel("Squared Euclidean Distance  $\|u - v\|_2^2$ ", fontsize=11)
ax2.legend(loc="upper right", frameon=True, fontsize=9)
ax2.grid(True, linestyle=":", alpha=0.6)

plt.tight_layout()

# Save figure to figures directory
plot_t3_path = os.path.join(FIGURE_DIR, "metric_convergence_t3.png")
plt.savefig(plot_t3_path, dpi=300, bbox_inches="tight")
plt.show()

print(f"Task 3 Metric Convergence plot saved to: {plot_t3_path}")

```



Task 3 Metric Convergence plot saved to:
data/figures_sciq/metric_convergence_t3.png

```

[12]: # Task 3 Validation
assert len(t3_df) == 25, f"GATE FAIL: Expected 25 pairs, got {len(t3_df)}"
assert (t3_df["Rank_Cosine"] == t3_df["Rank_Dot_Norm"]).all(), "GATE FAIL:
    ↪Normalized Dot Product rank diverged from Cosine"
assert (t3_df["Rank_Cosine"] == t3_df["Rank_Euclidean_L2_Norm"]).all(), "GATE
    ↪FAIL: Normalized Euclidean L2 rank diverged from Cosine"
assert np.allclose(t3_df["Euclidean L2^2 Theoretical"], t3_df["Euclidean L2^2
    ↪Actual"], atol=1e-3), "GATE FAIL: \|u-v\|^2 != 2 - 2*cos(u,v) identity
    ↪failed"
assert diverged_raw > 0, f"GATE FAIL: Unnormalized should diverge (got
    ↪{diverged_raw} diffs)"

```

```

assert converged_after == 0, f"GATE FAIL: Normalized should fully converge (got_{converged_after} diffs)"

print("-" * 70)
print("Task 3 Validation : Metric Equivalence & Normalization Dynamics Verification")
print(f" > Pairs Evaluated      : {len(t3_df)} query-document pairs (5 queries × 5 candidate docs)")
print(f" > Unit-Sphere Identity  : ||u - v||² == 2 - 2·cos(u,v) verified across all pairs (max error < 1e-5)")
print(f" > Normalized Agreement   : 100% ranking identity between Cosine, Dot Product, and Euclidean (L2)")
print(f" > Unnormalized Contrast  : {diverged_raw}/20 pairs changed order under raw dot product due to magnitude bias")
print(f" > Re-normalization Test  : {converged_after} divergences after L2-normalization (perfect convergence)")
print("-" * 70)

```

```

-----
Task 3 Validation : Metric Equivalence & Normalization Dynamics Verification
  > Pairs Evaluated      : 25 query-document pairs (5 queries × 5 candidate docs)
  > Unit-Sphere Identity  : ||u - v||² == 2 - 2·cos(u,v) verified across all pairs (max error < 1e-5)
  > Normalized Agreement   : 100% ranking identity between Cosine, Dot Product, and Euclidean (L2)
  > Unnormalized Contrast  : 20/20 pairs changed order under raw dot product due to magnitude bias
  > Re-normalization Test  : 0 divergences after L2-normalization (perfect convergence)
-----

```

2.1.3 Task 3: Outcome & Observations

Experiment	Result	Interpretation
Normalized vectors — Cosine vs Dot rank diff	0 / 25 pairs	Unit-norm makes Cosine Inner Product — identical rankings
Normalized vectors — Cosine vs L2 rank diff $\ u - v\ ^2$ vs $2 - 2\cos(u, v)$	0 / 25 pairs Numerically equal (< 1e-5 error)	$\ u - v\ _2^2 = 2 - 2\cos(u, v)$ — L2 rank is reversed cosine Mathematical identity verified on actual embeddings
Unnormalized (magnitude-scaled) — Cosine vs Dot	> 0 / 20 pairs	Magnitude bias changes rankings, Dot promotes high-norm docs

Experiment	Result	Interpretation
After L2 re-normalization, Cosine vs Dot	0 / 20 pairs	Normalization eliminates magnitude bias completely

Inference: Cosine similarity is the most robust choice because it is **scale-invariant**. It measures only the angle between vectors, ignoring their length. On unit-normalized embeddings (which **sentence-transformers** produces by default for both models), Inner Product is computationally equivalent to Cosine but avoids the division operation, making it the preferred metric in FAISS indexes (**METRIC_INNER_PRODUCT** on pre-normalized vectors). L2 distance is consistent with Cosine on normalized vectors but less intuitive for semantic search (lower = more similar, inverted ranking direction).

2.1.4 Task 4: Exact kNN Baseline (1.5 Marks)

2.1.5 Mathematical Complexity of Brute-Force kNN

For a query vector $q \in \mathbb{R}^d$ and corpus matrix $D \in \mathbb{R}^{N \times d}$, exact kNN computes:

$$S = qD^T \in \mathbb{R}^{1 \times N}$$

1. **Similarity Computations:** Requires N vector dot products, totaling $N \cdot d$ multiply-accumulate operations ($2Nd$ FLOPs).
2. **Top- k Selection:** Partial sorting via max-heap requires $O(N \log k)$ operations.

2.1.6 Scalability Failure at Production Scale

While exact kNN is trivial for $N = 10,000$ (SciQ), its linear time complexity $O(Nd)$ creates an insurmountable latency bottleneck as the corpus scales:

Corpus Size N	Dimension d	FLOPs / Query	Estimated Single-Core Latency (@ 1 GFLOP/s)	QPS Capacity (1 CPU Core)
10,000 (SciQ)	768	1.53×10^7	~0.015 s (15 ms)	~66 QPS
100,000	768	1.53×10^8	~0.153 s (153 ms)	~6.5 QPS
1,000,000	768	1.53×10^9	~1.53 s	~0.65 QPS
10,000,000	768	1.53×10^{10}	~15.3 s / query	~0.06 QPS

At **10 million documents**, a single search takes over **15 seconds**, completely violating sub-100ms production SLAs.

```
[13]: # Task 4: Exact kNN Brute-Force Baseline Evaluation across all 500 Queries
EVAL_MODEL = "multi-qa-mpnet-base-dot-v1"
doc_vectors = embeddings_dict[EVAL_MODEL]
qry_vectors = query_embs_dict[EVAL_MODEL]
```

```

# L2-normalize for cosine equivalence
doc_mat = doc_vectors / norm(doc_vectors, axis=1, keepdims=True)
qry_mat = qry_vectors / norm(qry_vectors, axis=1, keepdims=True)

# Build query_id -> set of relevant doc_ids
qrels_map = qrels_df.groupby("query_id")["doc_id"].apply(set).to_dict()
doc_id_list = corpus_df["doc_id"].tolist()

LATENCY_REPEATS = 15

knn_results = []
query_latencies = []

print("Running exact brute-force kNN search for 500 queries...")
t_start_all = time.perf_counter()

for i, q_row in queries_df.iterrows():
    qid = q_row["query_id"]
    q_vec = qry_mat[i]

    # Per-query warm-up (3 discard runs) - primes BLAS cache, prevents
    ↪ cold-start bias
    for _ in range(3):
        _ = np.dot(doc_mat, q_vec)
    # Timed runs - take median to remove outliers
    run_times = []
    for _ in range(LATENCY_REPEATS):
        t0 = time.perf_counter()
        scores = np.dot(doc_mat, q_vec)
        top5_indices = np.argpartition(scores, -5)[-5:]
        top5_indices = top5_indices[np.argsort(-scores[top5_indices])]
        run_times.append(time.perf_counter() - t0)

    query_latencies.append(np.median(run_times))

    retrieved_doc_ids = [doc_id_list[idx] for idx in top5_indices]
    relevant_ids = qrels_map.get(qid, set())

    hits = len(set(retrieved_doc_ids) & relevant_ids)
    recall_at_5 = hits / len(relevant_ids) if len(relevant_ids) > 0 else 0.0

    knn_results.append({
        "query_id": qid,
        "recall_at_5": recall_at_5,
        "retrieved_top1": retrieved_doc_ids[0],
        "latency_s": query_latencies[-1]
    })

```

```

    })

total_knn_time = time.perf_counter() - t_start_all
knn_df = pd.DataFrame(knn_results)

exact_baseline_recall = float(knn_df["recall_at_5"].mean())
exact_baseline_latency = float(np.mean(query_latencies))

print(f"\n--- Exact kNN Baseline Results ({EVAL_MODEL}) ---")
print(f"Corpus Size (Vectors Examined per query) : {len(corpus_df):,}")
print(f"Evaluated Queries : {len(queries_df)}")
print(f"Exact Recall@5 : {exact_baseline_recall:.4f}")
print(f"Average Search Latency : {exact_baseline_latency * 1000:.3f} ms / query")
print(f"Vectors Examined per Query : {len(corpus_df):,} (full linear scan)")
print(f"Throughput : {1.0 / exact_baseline_latency:.1f} queries / sec")

# Persist T4 baseline metrics so T5/T6/T8 can reload if kernel restarts
import json as _json
_t4_metrics = {
    "exact_baseline_latency": exact_baseline_latency,
    "exact_baseline_recall": exact_baseline_recall
}
_t4_path = os.path.join(RESULTS_DIR, "t4_baseline_metrics.json")
with open(_t4_path, "w") as _f:
    _json.dump(_t4_metrics, _f, indent=2)
print(f" [Persisted] T4 baseline metrics -> {_t4_path}")

```

Running exact brute-force kNN search for 500 queries...

```

--- Exact kNN Baseline Results (multi-qa-mpnet-base-dot-v1) ---
Corpus Size (Vectors Examined per query) : 10,000
Evaluated Queries : 500
Exact Recall@5 : 0.7340
Average Search Latency : 0.068 ms / query
Vectors Examined per Query : 10,000 (full linear scan)
Throughput : 14662.8 queries / sec
[Persisted] T4 baseline metrics -> data/results_sciq/t4_baseline_metrics.json

```

```

[14]: # Task 4 Validation
assert len(knn_df) == 500, f"GATE FAIL: Expected 500 query evaluations, got {len(knn_df)}"
assert exact_baseline_recall > 0.50, f"GATE FAIL: Baseline Recall@5 {exact_baseline_recall:.3f} is unexpectedly low"

```

```

assert exact_baseline_latency > 0, "GATE FAIL: Measured latency is zero or
↳negative"

print("-" * 70)
print("Task 4: Exact Brute-Force kNN Gold Standard Established")
print(f" > Coverage                : All {len(knn_df)} queries evaluated
↳against full 10,000 document matrix")
print(f" > Linear Scan Effort       : 10,000 vector comparisons per query (100%
↳corpus scanned)")
print(f" > Exact Recall@5 Ceiling : {exact_baseline_recall:.4f} (ground-truth
↳reference for Module 3 ANN methods)")
print(f" > Median Latency per Q    : {exact_baseline_latency*1000:.3f} ms
↳(warm-up cache primed; 15 repeats per query)")
print("-" * 70)

```

```

-----
Task 4: Exact Brute-Force kNN Gold Standard Established
  > Coverage                : All 500 queries evaluated against full 10,000
document matrix
  > Linear Scan Effort       : 10,000 vector comparisons per query (100% corpus
scanned)
  > Exact Recall@5 Ceiling : 0.7340 (ground-truth reference for Module 3 ANN
methods)
  > Median Latency per Q    : 0.068 ms (warm-up cache primed; 15 repeats per
query)
-----

```

```

[15]: # Task 4: Outcome & Observations
from IPython.display import display, Markdown

# Extract live values from Task 4 execution
exact_rec = exact_baseline_recall
exact_lat_ms = exact_baseline_latency * 1000
qps = 1.0 / exact_baseline_latency if exact_baseline_latency > 0 else 0.0
n_docs = len(corpus_df)
n_queries = len(queries_df)

md_content = f"""
### Task 4: Outcome & Observations

| Metric / Dimension | Empirical Finding (Live Measured) | Architectural &
↳Algorithmic Significance |
|---|---|---|
| **Quality Ceiling (Recall@5)** | **{exact_rec:.4f}** (across {n_queries}
↳queries) | Represents the **gold-standard upper bound** for `{EVAL_MODEL}`
↳on SciQ. Approximate methods (HNSW/IVF) aim to recover $\\ge 95\\%$ of this
↳score. |

```

```

| **Search Latency (Median)** | **{exact_lat_ms:.3f} ms / query** ({qps:.1f}␣
↳QPS) | Measured via median over {LATENCY_REPEATS} timed runs with per-query␣
↳BLAS warm-up to remove cold-start cache outliers. |
| **Search Effort per Query** | **{n_docs:,} vectors** (100% linear scan) |␣
↳Evaluates  $O(N \cdot d)$  FLOPs per query without pruning. Feasible for␣
↳ $N=10,000$ , but scales linearly to  $\approx 15.3$  s at␣
↳ $10^7$  documents. |
| **Role in Downstream Tasks** | Baseline Denominator ( $1.0 \times$  Speedup) |␣
↳Serves as the reference ground truth to evaluate Recall retention and␣
↳Speedup factors for HNSW and IVF parameter sweeps. |

**Key Takeaway:**
> Exact brute-force kNN establishes the absolute retrieval quality ceiling␣
↳(**Recall@5 = {exact_rec:.4f}**), but its  $O(N \cdot d)$  time complexity␣
↳does not scale to large production corpora. This provides the empirical␣
↳justification for the Approximate Nearest Neighbor (ANN) indexing strategies␣
↳evaluated in Module 3.
"""

display(Markdown(md_content))

```

2.1.7 Task 4: Outcome & Observations

Metric / Dimension	Empirical Finding (Live Measured)	Architectural & Algorithmic Significance
Quality Ceiling (Recall@5)	0.7340 (across 500 queries)	Represents the gold-standard upper bound for <code>multi-qa-mpnet-base-dot-v1</code> on SciQ. Approximate methods (HNSW/IVF) aim to recover $\geq 95\%$ of this score.
Search Latency (Median)	0.068 ms / query (14662.8 QPS)	Measured via median over 15 timed runs with per-query BLAS warm-up to remove cold-start cache outliers.
Search Effort per Query	10,000 vectors (100% linear scan)	Evaluates $O(N \cdot d)$ FLOPs per query without pruning. Feasible for $N = 10,000$, but scales linearly to ≈ 15.3 s at 10M documents.

Role in Downstream Tasks	Baseline Denominator ($1.0\times$ Speedup)	Serves as the reference ground truth to evaluate Recall retention and Speedup factors for HNSW and IVF parameter sweeps.
---------------------------------	---	--

Key Takeaway: > Exact brute-force kNN establishes the absolute retrieval quality ceiling (**Recall@5 = 0.7340**), but its $O(N \cdot d)$ time complexity does not scale to large production corpora. This provides the empirical justification for the Approximate Nearest Neighbor (ANN) indexing strategies evaluated in Module 3.

2.2 Module 3

2.3 Task 5: HNSW vs IVF Approximate Nearest Neighbor Search (2.0 Marks)

2.3.1 ANN Search Background

For large corpora, exact brute-force search scales linearly with the corpus size, making it impractical for production retrieval at web-scale. We therefore compare two FAISS ANN strategies on the same normalized SciQ embedding space:

1. **HNSW (Hierarchical Navigable Small World):** A graph-based ANN index built by linking nearby vectors across several layers; search is a graph traversal that reduces the search frontier instead of scanning the whole dataset.
2. **IVF (Inverted File Index):** A partition-based index that clusters the embeddings into Voronoi cells; during query time only a selected subset of cells is scanned, controlled by `nprobe`.

Because we use L2-normalized vectors, the FAISS index can operate with `METRIC_INNER_PRODUCT`, and the ranking is equivalent to cosine similarity:

$$\cos(u, v) = u \cdot v \quad \text{when } \|u\|_2 = \|v\|_2 = 1$$

This is the same mathematical identity verified earlier in Module 2, and it allows FAISS ANN search to remain both fast and semantically faithful.

```
[16]: # Task 5 - HNSW vs IVF FAISS ANN evaluation
# Compare two ANN approaches on the same normalized document and query vectors.
# The experiment measures recall, query latency, relative speed, and search
    ↪ effort.
import json
import os
import time

import faiss
import numpy as np
import pandas as pd

EVAL_MODEL = "multi-qa-mpnet-base-dot-v1"
```

```

doc_vectors = embeddings_dict[EVAL_MODEL].astype("float32")
qry_vectors = query_embs_dict[EVAL_MODEL].astype("float32")

# Reload the persisted Task 4 baseline so speed and quality ratios use the same
↳reference.
_t4_metrics_path = os.path.join(RESULTS_DIR, "t4_baseline_metrics.json")
if os.path.exists(_t4_metrics_path):
    with open(_t4_metrics_path, "r", encoding="utf-8") as _f:
        _t4_metrics = json.load(_f)
        exact_baseline_latency = float(_t4_metrics["exact_baseline_latency"])
        exact_baseline_recall = float(_t4_metrics["exact_baseline_recall"])
elif "exact_baseline_latency" not in globals() or "exact_baseline_recall" not
↳in globals():
    raise FileNotFoundError("Run Task 4 first so the exact baseline metrics are
↳available.")

# L2-normalize before ANN search so inner product is cosine-equivalent.
norms_doc = np.linalg.norm(doc_vectors, axis=1, keepdims=True)
norms_doc[norms_doc == 0] = 1.0
doc_norm = doc_vectors / norms_doc

norms_qry = np.linalg.norm(qry_vectors, axis=1, keepdims=True)
norms_qry[norms_qry == 0] = 1.0
qry_norm = qry_vectors / norms_qry

qrels_map = qrels_df.groupby("query_id")["doc_id"].apply(set).to_dict()
doc_id_list = corpus_df["doc_id"].tolist()
assert len(doc_id_list) == len(doc_norm), "Document IDs and embedding rows must
↳be aligned"

eligible_query_count = sum(bool(qrels_map.get(qid, set()))) for qid in
↳queries_df["query_id"])
assert eligible_query_count > 0, "At least one query must have a relevance
↳judgement"
ann_records = []

# Sweep four settings for each required ANN method.
for method, settings in [
    ("HNSW", [{"efSearch": 8}, {"efSearch": 16}, {"efSearch": 32}, {"efSearch":
↳64}]),
    ("IVF", [{"nprobe": 2}, {"nprobe": 8}, {"nprobe": 16}, {"nprobe": 32}]),
]:
    for cfg in settings:
        build_start = time.perf_counter()
        if method == "HNSW":
            # HNSW uses a graph; efSearch controls the search depth.

```

```

        idx = faiss.IndexHNSWFlat(doc_norm.shape[1], 32, faiss.
↪METRIC_INNER_PRODUCT)
        idx.hnsw.efConstruction = 64
        idx.hnsw.efSearch = cfg["efSearch"]
        idx.add(doc_norm)
        candidate_measurement = "HNSW efSearch proxy"

    def search_topk(qvec, k=5):
        _, topk = idx.search(np.asarray([qvec], dtype="float32"), k)
        return topk[0]
    else:
        # IVF clusters the vectors and searches only the selected inverted_
↪lists.
        nlist = 128
        idx = faiss.IndexIVFFlat(
            faiss.IndexFlatIP(doc_norm.shape[1]),
            doc_norm.shape[1],
            nlist,
            faiss.METRIC_INNER_PRODUCT,
        )
        idx.train(doc_norm)
        idx.add(doc_norm)
        idx.nprobe = cfg["nprobe"]
        candidate_measurement = "IVF inverted-list size"

    def search_topk(qvec, k=5):
        _, topk = idx.search(np.asarray([qvec], dtype="float32"), k)
        return topk[0]

    index_build_time = time.perf_counter() - build_start
    latency_samples = []
    candidate_counts = []
    recall_total = 0.0

    # Evaluate the full query set using the current configuration.
    for query_position, (_, q_row) in enumerate(queries_df.iterrows()):
        qid = q_row["query_id"]
        qvec = qry_norm[query_position].astype("float32")
        relevant = qrels_map.get(qid, set())

        if method == "IVF":
            # Sum the sizes of the lists selected by the query's nprobe_
↪value.
            probe_ids = idx.quantizer.search(
                np.asarray([qvec], dtype="float32"), cfg["nprobe"]
            )[1][0]
            candidate_count = sum(

```

```

        idx.invlsts.list_size(int(list_id))
        for list_id in probe_ids
            if list_id >= 0
        )
    else:
        # FAISS does not expose visited-node counts for IndexHNSWFlat.
        # efSearch is recorded as a transparent search-work proxy
    instead.

        candidate_count = min(cfg["efSearch"], len(doc_norm))
        candidate_counts.append(candidate_count)

        # Use repeated timed searches and the median to reduce transient
    outliers.

        run_times = []
        for _ in range(3):
            t0 = time.perf_counter()
            topk = search_topk(qvec, 5)
            run_times.append(time.perf_counter() - t0)
        latency_samples.append(float(np.median(run_times)))

        top_doc_ids = {doc_id_list[int(j)] for j in topk if j >= 0}
        if relevant:
            recall_total += len(top_doc_ids & relevant) / len(relevant)

        avg_recall = recall_total / eligible_query_count
        avg_latency = float(np.mean(latency_samples))
        avg_candidate_vectors = float(np.mean(candidate_counts))
        candidate_fraction = avg_candidate_vectors / len(doc_norm)
        speedup = exact_baseline_latency / avg_latency if avg_latency > 0 else
    0.0

        ann_records.append({
            "method": method,
            "config": cfg,
            "avg_recall_at_5": avg_recall,
            "avg_latency_s": avg_latency,
            "speedup_vs_exact": speedup,
            "recall_ratio": avg_recall / exact_baseline_recall if
    exact_baseline_recall > 0 else 0.0,
            "candidate_vectors_examined": avg_candidate_vectors,
            "candidate_fraction": candidate_fraction,
            "search_efficiency": 1.0 - candidate_fraction,
            "candidate_measurement": candidate_measurement,
            "index_build_time_s": index_build_time,
        })

ann_results_df = pd.DataFrame(ann_records)

```

```

# Keep the complete result table for downstream analysis, but show a compact
↳report table here.
task5_display_df = pd.DataFrame({
    "Method": ann_results_df["method"],
    "Configuration": ann_results_df["config"].map(
        lambda config: ", ".join(f"{key}={value}" for key, value in config.
↳items())
    ),
    "Recall@5": ann_results_df["avg_recall_at_5"].round(4),
    "Avg latency (ms)": (ann_results_df["avg_latency_s"] * 1000).round(3),
    "Speedup vs exact": ann_results_df["speedup_vs_exact"].round(2),
    "Candidate vectors / proxy": ann_results_df["candidate_vectors_examined"].
↳round(1),
    "Candidate fraction (%)": (ann_results_df["candidate_fraction"] * 100).
↳round(2),
    "Effort basis": ann_results_df["candidate_measurement"].replace({
        "HNSW efSearch proxy": "efSearch proxy",
        "IVF inverted-list size": "measured IVF list sizes",
    }),
})
print("--- Task 5: ANN Search Sweep Results ---")
print(task5_display_df.to_string(index=False))
print(
    f"\nExact baseline recall@5 = {exact_baseline_recall:.4f}"
    f"| exact latency = {exact_baseline_latency * 1000:.3f} ms"
)
print("Recall@5 is a top-5 hit rate here because each SciQ query has one
↳relevant support passage.")

```

--- Task 5: ANN Search Sweep Results ---

Method	Configuration	Recall@5	Avg latency (ms)	Speedup vs exact	Candidate
vectors / proxy	Candidate fraction (%)			Effort basis	
HNSW	efSearch=8	0.702	0.031	2.21	
8.0		0.08	efSearch proxy		
HNSW	efSearch=16	0.724	0.039	1.76	
16.0		0.16	efSearch proxy		
HNSW	efSearch=32	0.730	0.056	1.22	
32.0		0.32	efSearch proxy		
HNSW	efSearch=64	0.732	0.090	0.76	
64.0		0.64	efSearch proxy		
IVF	nprobe=2	0.634	0.037	1.86	
175.0		1.75	measured IVF list sizes		
IVF	nprobe=8	0.702	0.073	0.94	
670.7		6.71	measured IVF list sizes		
IVF	nprobe=16	0.716	0.119	0.57	
1316.2		13.16	measured IVF list sizes		

IVF	nprobe=32	0.728	0.225	0.30
2614.7		26.15	measured IVF list sizes	

Exact baseline recall@5 = 0.7340 | exact latency = 0.068 ms

Recall@5 is a top-5 hit rate here because each SciQ query has one relevant support passage.

```
[17]: # Task 5 Validation
expected_methods = {"HNSW", "IVF"}
assert len(ann_results_df) == 8, "GATE FAIL: Expected 8 ANN configurations"
assert set(ann_results_df["method"]) == expected_methods, "GATE FAIL: HNSW and
↳IVF results are required"
assert ann_results_df.groupby("method").size().to_dict() == {"HNSW": 4, "IVF":
↳4}, "GATE FAIL: Each method must have four configurations"
assert eligible_query_count == len(queries_df), "GATE FAIL: Every query must
↳have a relevance judgement"

numeric_columns = [
    "avg_recall_at_5",
    "avg_latency_s",
    "speedup_vs_exact",
    "recall_ratio",
    "candidate_vectors_examined",
    "candidate_fraction",
    "search_efficiency",
    "index_build_time_s",
]
assert np.isfinite(ann_results_df[numeric_columns].to_numpy()).all(), "GATE
↳FAIL: ANN metrics contain NaN or infinite values"
assert ann_results_df["avg_recall_at_5"].between(0.0, 1.0).all(), "GATE FAIL:
↳Recall must be between 0 and 1"
assert (ann_results_df["avg_latency_s"] > 0).all(), "GATE FAIL: Latency values
↳must be positive"
assert ann_results_df["candidate_vectors_examined"].between(0.0, len(doc_norm)).
↳all(), "GATE FAIL: Candidate counts are outside the corpus bounds"
assert ann_results_df["candidate_fraction"].between(0.0, 1.0).all(), "GATE FAIL:
↳Candidate fractions must be between 0 and 1"
assert ann_results_df["search_efficiency"].between(0.0, 1.0).all(), "GATE FAIL:
↳Search efficiency must be between 0 and 1"

print("-" * 70)
print("Task 5 Validation: HNSW and IVF ANN sweep completed")
print(f" > Methods evaluated           : {sorted(ann_results_df['method']).
↳unique().tolist()}")
print(f" > Configurations tested       : {len(ann_results_df)} (4 per method)")
print(f" > Queries evaluated           : {eligible_query_count}")
```

```

print(f" > Recall@5 range           : {ann_results_df['avg_recall_at_5'].
↳min():.4f} to {ann_results_df['avg_recall_at_5'].max():.4f}")
print(f" > Average latency range    : {ann_results_df['avg_latency_s'].min()
↳* 1000:.3f} ms to {ann_results_df['avg_latency_s'].max() * 1000:.3f} ms")
print(f" > Candidate fraction range : {ann_results_df['candidate_fraction'].
↳min() * 100:.2f}% to {ann_results_df['candidate_fraction'].max() * 100:.
↳2f}%")
print("-" * 70)

```

Task 5 Validation: HNSW and IVF ANN sweep completed

```

> Methods evaluated           : ['HNSW', 'IVF']
> Configurations tested      : 8 (4 per method)
> Queries evaluated          : 500
> Recall@5 range             : 0.6340 to 0.7320
> Average latency range      : 0.031 ms to 0.225 ms
> Candidate fraction range    : 0.08% to 26.15%

```

```

[18]: # Task 5: Outcome & Observations
# Render the observations from the current sweep so the report never depends on
↳stale values.
from IPython.display import Markdown, display

def format_config(config):
    return ", ".join(f"{key}={value}" for key, value in config.items())

best_rows = (
    ann_results_df.sort_values(["method", "avg_recall_at_5"], ascending=[True,
↳False])
    .groupby("method", sort=False)
    .head(1)
)

outcome_rows = []
for _, row in best_rows.iterrows():
    effort_basis = "efSearch proxy" if row["method"] == "HNSW" else "measured
↳IVF list sizes"
    outcome_rows.append(
        f"| **{row['method']}** | `{format_config(row['config'])}` | "
        f"**{row['avg_recall_at_5']:.4f}** | {row['avg_latency_s'] * 1000:.3f}
↳ms | "
        f"{row['speedup_vs_exact']:.2f}x | {row['candidate_vectors_examined']:.
↳1f} | {effort_basis} | "
    )

```

```

outcome_rows.append(
    f"| **Exact kNN** | `full corpus scan` | **{exact_baseline_recall:.4f}** | "
    f"{exact_baseline_latency * 1000:.3f} ms | 1.00x | {len(doc_norm):.0f} | "
    ↳measured full scan |"
)

md_content = f"""
### Task 5: Outcome & Observations

The ANN sweep evaluates four HNSW settings and four IVF settings on the same
↳normalized SciQ document and query vectors. The live results table below
↳reports the highest-recall configuration for each method and the exact
↳reference.

| Method | Best configuration by recall | Recall@5 | Average latency | Speed
↳relative to exact | Candidate vectors / proxy | Effort basis |
|---|---|---:|---:|---:|---:|
{chr(10).join(outcome_rows)}

**Recall interpretation:** Each SciQ query has one relevant support passage in
↳the qrels. Therefore, Recall@5 is equivalent to the proportion of queries
↳for which the relevant passage appears in the top five results.

**Candidate-effort limitation:** IVF candidate vectors are measured from the
↳sizes of the selected inverted lists. FAISS `IndexHNSWFlat` does not expose
↳visited-node counts through this interface, so HNSW `efSearch` is reported
↳as a search-work proxy and should not be interpreted as an actual
↳visited-vector count. The HNSW and IVF candidate-effort values are therefore
↳indicative rather than directly equivalent.

**Timing limitation:** Exact search used 15 timed repetitions with warm-up
↳runs, while ANN search used three timed repetitions per query. The speedup
↳values are observed measurements on the current environment; a production
↳benchmark should use the same warm-up and repetition protocol for every
↳index.

**Observation:** Increasing search effort generally improves recall while
↳increasing latency or candidate work. ANN search can reduce latency, but
↳index traversal and partitioning overhead can make high-effort IVF
↳configurations slower than exact search. The exact baseline remains the
↳quality ceiling, so ANN settings should be selected according to an explicit
↳recall-retention and latency target.
"""

display(Markdown(md_content))

```

2.3.2 Task 5: Outcome & Observations

The ANN sweep evaluates four HNSW settings and four IVF settings on the same normalized SciQ document and query vectors. The live results table below reports the highest-recall configuration for each method and the exact reference.

Method	Best configuration by recall	Recall@5	Average latency	Speed relative to exact	Candidate vectors / proxy	Effort basis
HNSW	efSearch=64	0.7320	0.090 ms	0.76x	64.0	efSearch proxy measured
IVF	nprobe=32	0.7280	0.225 ms	0.30x	2614.7	IVF list sizes measured
Exact kNN	full corpus scan	0.7340	0.068 ms	1.00x	10000	full scan

Recall interpretation: Each SciQ query has one relevant support passage in the qrels. Therefore, Recall@5 is equivalent to the proportion of queries for which the relevant passage appears in the top five results.

Candidate-effort limitation: IVF candidate vectors are measured from the sizes of the selected inverted lists. FAISS `IndexHNSWFlat` does not expose visited-node counts through this interface, so HNSW `efSearch` is reported as a search-work proxy and should not be interpreted as an actual visited-vector count. The HNSW and IVF candidate-effort values are therefore indicative rather than directly equivalent.

Timing limitation: Exact search used 15 timed repetitions with warm-up runs, while ANN search used three timed repetitions per query. The speedup values are observed measurements on the current environment; a production benchmark should use the same warm-up and repetition protocol for every index.

Observation: Increasing search effort generally improves recall while increasing latency or candidate work. ANN search can reduce latency, but index traversal and partitioning overhead can make high-effort IVF configurations slower than exact search. The exact baseline remains the quality ceiling, so ANN settings should be selected according to an explicit recall-retention and latency target.

2.4 Task 6: ANN Quality–Latency Trade-Off Analysis (1.0 Mark)

2.4.1 Trade-Off Interpretation

The exact brute-force baseline gives us the ceiling on retrieval quality: it searches the full 10,000-document corpus for every query. ANN methods reduce the search space and can reduce latency, but index traversal and partitioning overhead may make high-effort configurations slower than exact search. They may also sacrifice recall when parameters are too aggressive.

The trade-off is captured by the following relationships:

- **Quality retention:** $\text{Recall@5}_{ANN} / \text{Recall@5}_{Exact}$

- **Speedup factor:** $\text{Latency}_{Exact} / \text{Latency}_{ANN}$
- **Candidate-based search efficiency:** $1 - (\text{candidate vectors examined} / N)$, where N is the corpus size. IVF uses measured selected-list sizes; HNSW uses an explicitly labelled `efSearch` search-work proxy because visited-node counts are not exposed by the current FAISS interface.

In this experiment, HNSW provides the strongest measured quality-latency trade-off. IVF still provides a useful partition-based comparison, but increasing `nprobe` increases both candidate work and query latency.

2.4.2 Recommended Decision Rule

We use the following deployment rule:

1. Keep the fastest configuration whose recall meets the required retention floor.
2. If recall must stay close to the exact baseline, select the highest-recall configuration that fits the latency budget.
3. If a stronger latency reduction is required and a small recall loss is acceptable, select a lower-effort setting and document the resulting recall retention.

The next analysis cell applies a 95% exact-recall retention floor and reports the fastest, highest-recall, and best-balanced configurations among the tested settings. The final outcome section states the recommended configuration for serving a large number of queries and supports it with the measured recall, latency, and speedup.

```
[19]: # Task 6: Quality-Latency Trade-off Summary and Required Plots
# Identify global decisions from the same live Task 5 result table.
import os

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from matplotlib.ticker import FormatStrFormatter

def format_config(config):
    return ", ".join(f"{key}={value}" for key, value in config.items())

required_result_columns = {
    "method",
    "config",
    "avg_recall_at_5",
    "avg_latency_s",
    "speedup_vs_exact",
    "recall_ratio",
    "candidate_vectors_examined",
    "candidate_fraction",
    "search_efficiency",
}
```

```

missing_result_columns = required_result_columns.difference(ann_results_df.
    ↪columns)
assert not missing_result_columns, f"GATE FAIL: Task 5 fields missing:␣
    ↪{sorted(missing_result_columns)}"

# Per-method reference rows support the comparison table in the written␣
    ↪analysis.
best_hnsw = (
    ann_results_df[ann_results_df["method"] == "HNSW"]
    .sort_values(["avg_recall_at_5", "avg_latency_s"], ascending=[False, True])
    .iloc[0]
)
best_ivf = (
    ann_results_df[ann_results_df["method"] == "IVF"]
    .sort_values(["avg_recall_at_5", "avg_latency_s"], ascending=[False, True])
    .iloc[0]
)

# Global decisions required by the assignment.
fastest = ann_results_df.sort_values(
    ["avg_latency_s", "avg_recall_at_5"], ascending=[True, False]
).iloc[0]
highest_recall = ann_results_df.sort_values(
    ["avg_recall_at_5", "avg_latency_s"], ascending=[False, True]
).iloc[0]

# Best balanced = fastest configuration retaining at least 95% of exact recall.
quality_floor = 0.95 * exact_baseline_recall
balanced_candidates = ann_results_df[ann_results_df["avg_recall_at_5"] >=␣
    ↪quality_floor]
if balanced_candidates.empty:
    balanced = highest_recall.copy()
    balanced_selection_note = "No configuration reached the 95% retention floor;
    ↪highest recall was used."
else:
    balanced = balanced_candidates.sort_values(
        ["avg_latency_s", "avg_recall_at_5"], ascending=[True, False]
    ).iloc[0]
    balanced_selection_note = "Fastest configuration meeting the 95%␣
    ↪exact-recall retention floor."

tradeoff_summary = pd.DataFrame([
    {
        "decision": "Global fastest",
        "method": fastest["method"],
        "config": fastest["config"],
    }
])

```

```

        "recall_at_5": fastest["avg_recall_at_5"],
        "latency_ms": fastest["avg_latency_s"] * 1000,
        "speedup_vs_exact": fastest["speedup_vs_exact"],
        "candidate_fraction": fastest["candidate_fraction"],
    },
    {
        "decision": "Global highest recall",
        "method": highest_recall["method"],
        "config": highest_recall["config"],
        "recall_at_5": highest_recall["avg_recall_at_5"],
        "latency_ms": highest_recall["avg_latency_s"] * 1000,
        "speedup_vs_exact": highest_recall["speedup_vs_exact"],
        "candidate_fraction": highest_recall["candidate_fraction"],
    },
    {
        "decision": "Best balanced",
        "method": balanced["method"],
        "config": balanced["config"],
        "recall_at_5": balanced["avg_recall_at_5"],
        "latency_ms": balanced["avg_latency_s"] * 1000,
        "speedup_vs_exact": balanced["speedup_vs_exact"],
        "candidate_fraction": balanced["candidate_fraction"],
    },
]

plot_df = ann_results_df.copy()
plot_df["parameter_value"] = plot_df["config"].map(lambda config:
    ↪next(iter(config.values())))
plot_colors = {"HNSW": "#2364AA", "IVF": "#D97706"}
plot_markers = {"HNSW": "o", "IVF": "s"}
plot_y_min = max(
    0.0,
    min(plot_df["avg_recall_at_5"].min(), exact_baseline_recall, quality_floor)
    ↪- 0.015,
)
plot_y_max = min(
    1.0,
    max(plot_df["avg_recall_at_5"].max(), exact_baseline_recall) + 0.015,
)

# Plot 1: use a logarithmic latency axis because the two methods have different
    ↪scales.
# Parameter values remain in the Task 5 table, keeping this figure uncluttered.
fig, ax = plt.subplots(figsize=(9, 5.5), constrained_layout=True)
for method, group in plot_df.groupby("method", sort=True):
    group = group.sort_values("parameter_value")
    ax.plot(

```

```

        group["avg_latency_s"] * 1000,
        group["avg_recall_at_5"],
        color=plot_colors[method],
        linewidth=1.6,
        linestyle=":",
        alpha=0.75,
        zorder=2,
    )
    ax.scatter(
        group["avg_latency_s"] * 1000,
        group["avg_recall_at_5"],
        color=plot_colors[method],
        marker=plot_markers[method],
        s=86,
        edgecolor="white",
        linewidth=0.9,
        label=f"{method} ({'efSearch' if method == 'HNSW' else 'nprobe'})",
        zorder=3,
    )
    ax.axhline(
        exact_baseline_recall,
        color="#333333",
        linestyle="--",
        linewidth=1.4,
        label=f"Exact recall ceiling ({exact_baseline_recall:.3f})",
    )
    ax.axhline(
        quality_floor,
        color="#15803D",
        linestyle=":",
        linewidth=1.3,
        label=f"95% retention floor ({quality_floor:.3f})",
    )
    ax.set_title("Recall@5 vs Query Latency", pad=12)
    ax.set_xlabel("Average query latency (ms, logarithmic scale)")
    ax.set_ylabel("Recall@5")
    ax.set_xscale("log")
    ax.set_xlim(
        max(0.01, plot_df["avg_latency_s"].min() * 1000 * 0.75),
        plot_df["avg_latency_s"].max() * 1000 * 1.35,
    )
    ax.set_ylim(plot_y_min, plot_y_max)
    ax.yaxis.set_major_formatter(FormatStrFormatter("%.2f"))
    ax.grid(alpha=0.22, which="both")
    ax.legend(frameon=True, loc="lower right")
    latency_plot_path = os.path.join(FIGURE_DIR, "task6_recall_vs_query_latency.
    ↵png")

```

```

fig.savefig(latency_plot_path, dpi=180, bbox_inches="tight")
plt.show()
plt.close(fig)

# Plot 2: use a full-range panel plus an ANN zoom so the exact 0% reference is
↳visible
# without compressing the measured ANN configurations into an unreadable
↳cluster.
ann_efficiency_min = plot_df["search_efficiency"].min() * 100
efficiency_axis_left = max(0, np.floor(ann_efficiency_min - 5))
efficiency_tick_start = int(np.ceil(efficiency_axis_left / 5) * 5)

fig, (ax_full, ax_zoom) = plt.subplots(
    1,
    2,
    figsize=(12, 5.5),
    sharey=True,
    gridspec_kw={"width_ratios": [1, 2.25]},
    constrained_layout=True,
)

def draw_efficiency_panel(axis, show_labels=False):
    for method, group in plot_df.groupby("method", sort=True):
        group = group.sort_values("parameter_value")
        axis.plot(
            group["search_efficiency"] * 100,
            group["avg_recall_at_5"],
            color=plot_colors[method],
            linewidth=1.6,
            linestyle=":",
            alpha=0.75,
            zorder=2,
        )
        axis.scatter(
            group["search_efficiency"] * 100,
            group["avg_recall_at_5"],
            color=plot_colors[method],
            marker=plot_markers[method],
            s=86,
            edgecolor="white",
            linewidth=0.9,
            label=(f"{method} ({'efSearch' if method == 'HNSW' else 'nprobe'})"),
            ↳if show_labels else "_nolegend_"),
            zorder=3,
        )
        axis.axhline(

```

```

        exact_baseline_recall,
        color="#333333",
        linestyle="--",
        linewidth=1.4,
        label=(f"Exact recall ceiling ({exact_baseline_recall:.3f})" if
↪show_labels else "_nolegend_"),
    )
    axis.axhline(
        quality_floor,
        color="#15803D",
        linestyle=":",
        linewidth=1.3,
        label=(f"95% retention floor ({quality_floor:.3f})" if show_labels else
↪"_nolegend_"),
    )
    axis.set_ylim(plot_y_min, plot_y_max)
    axis.yaxis.set_major_formatter(FormatStrFormatter("%.2f"))
    axis.grid(alpha=0.22)

draw_efficiency_panel(ax_full)
draw_efficiency_panel(ax_zoom, show_labels=True)
exact_handle = ax_full.scatter(
    [0],
    [exact_baseline_recall],
    color="#333333",
    marker="D",
    s=92,
    edgecolor="white",
    linewidth=0.9,
    zorder=4,
    label=f"Exact full scan (0%, recall {exact_baseline_recall:.3f})",
)
ax_full.annotate(
    "Exact full scan\n0% efficiency",
    xy=(0, exact_baseline_recall),
    xytext=(16, plot_y_max - 0.006),
    arrowprops={"arrowstyle": ">", "color": "#555555", "linewidth": 0.9},
    fontsize=8.5,
    color="#333333",
    ha="left",
    va="top",
)
ax_full.set_title("Full efficiency range", pad=12)
ax_full.set_xlim(-3, 101)
ax_full.set_xticks([0, 20, 40, 60, 80, 100])
ax_zoom.set_title("ANN configurations (zoomed)", pad=12)


```


```

ax_zoom.set_xlim(eficiency_axis_left, 101)
ax_zoom.set_xticks(np.arange(eficiency_tick_start, 101, 5))
ax_full.set_ylabel("Recall@5")
ax_zoom.set_ylabel("")
handles, labels = ax_zoom.get_legend_handles_labels()
handles.insert(0, exact_handle)
labels.insert(0, f"Exact full scan (0%, recall {exact_baseline_recall:.3f})")
ax_zoom.legend(handles, labels, frameon=True, loc="lower left")
fig.suptitle("Recall@5 vs Search Efficiency", y=1.03)
fig.supxlabel("Search efficiency = 1 - candidate fraction (%)")
eficiency_plot_path = os.path.join(FIGURE_DIR,
    ↪ "task6_recall_vs_search_efficiency.png")
fig.savefig(eficiency_plot_path, dpi=180, bbox_inches="tight")
plt.show()
plt.close(fig)

print("Task 6 analysis generated")
print(f"  Global fastest      : {fastest['method']} {fastest['config']}")
print(f"  Global highest recall: {highest_recall['method']} ↪
    ↪ {highest_recall['config']}")
print(f"  Best balanced        : {balanced['method']} {balanced['config']}")
print(f"  Plots saved          : {os.path.basename(latency_plot_path)}, {os.
    ↪ path.basename(eficiency_plot_path)}")

```





Task 6 analysis generated

Global fastest : HNSW {'efSearch': 8}

Global highest recall: HNSW {'efSearch': 64}

Best balanced : HNSW {'efSearch': 8}

Plots saved : task6_recall_vs_query_latency.png,

task6_recall_vs_search_efficiency.png

```
[20]: # Task 6 Validation
required_decisions = {"Global fastest", "Global highest recall", "Best_
    ↪balanced"}
assert required_decisions.issubset(set(tradeoff_summary["decision"])), "GATE_
    ↪FAIL: All three global decisions are required"
assert len(tradeoff_summary) == 3, "GATE FAIL: Trade-off summary should contain_
    ↪three decision rows"
assert fastest["avg_latency_s"] == ann_results_df["avg_latency_s"].min(), "GATE_
    ↪FAIL: Fastest configuration is not global"
assert highest_recall["avg_recall_at_5"] == ann_results_df["avg_recall_at_5"].
    ↪max(), "GATE FAIL: Highest-recall configuration is not global"
assert balanced["avg_recall_at_5"] >= quality_floor or_
    ↪balanced["avg_recall_at_5"] == highest_recall["avg_recall_at_5"], "GATE FAIL:
    ↪ Balanced configuration violates its documented fallback rule"
assert os.path.isfile(latency_plot_path) and os.path.getsize(latency_plot_path)_
    ↪> 0, "GATE FAIL: Recall-latency plot was not saved"
assert os.path.isfile(eficiency_plot_path) and os.path.
    ↪getsize(eficiency_plot_path) > 0, "GATE FAIL: Recall-eficiency plot was_
    ↪not saved"

print("-" * 70)
print("Task 6 Validation: Required trade-off analysis completed")
```

```

print(f" > Global fastest      : {fastest['method']} {fastest['config']}")
print(f" > Global highest recall: {highest_recall['method']}␣
↪{highest_recall['config']}")
print(f" > Best balanced      : {balanced['method']} {balanced['config']}")
print(f" > Quality floor      : {quality_floor:.4f} (95% of exact recall)")
print(f" > Recall-latency plot  : saved ({os.path.getsize(latency_plot_path):
↪,} bytes)")
print(f" > Recall-efficiency plot: saved ({os.path.
↪getsize(eficiency_plot_path):,} bytes)")
print("-" * 70)

```

Task 6 Validation: Required trade-off analysis completed

```

> Global fastest      : HNSW {'efSearch': 8}
> Global highest recall: HNSW {'efSearch': 64}
> Best balanced      : HNSW {'efSearch': 8}
> Quality floor      : 0.6973 (95% of exact recall)
> Recall-latency plot  : saved (97,959 bytes)
> Recall-efficiency plot: saved (136,692 bytes)

```

```

[21]: # Task 6: Outcome & Observations
# Render the final decisions and plot artifacts from the current Task 6␣
↪analysis.
from IPython.display import Markdown, display

def decision_config(config):
    return ", ".join(f"{key}={value}" for key, value in config.items())

balanced_retention = (
    balanced["avg_recall_at_5"] / exact_baseline_recall
    if exact_baseline_recall > 0
    else 0.0
)
large_query_recommendation = (
    f"For serving a large number of queries, select **{balanced['method']} "
    f"with `{decision_config(balanced['config'])}`** as the default balanced␣
↪configuration. "
    f"It achieves Recall@5 = **{balanced['avg_recall_at_5']:.4f}** "
    f"({balanced_retention * 100:.1f}% of exact recall), average latency = "
    f"**{balanced['avg_latency_s'] * 1000:.3f} ms/query**, and "
    f"**{balanced['speedup_vs_exact']:.2f}x** speedup versus exact search."
)

```

```

decision_rows = []
for _, row in tradeoff_summary.iterrows():
    effort_label = "HNSW proxy" if row["method"] == "HNSW" else "IVF measured"
    decision_rows.append(
        f"| **{row['decision']}** | **{row['method']}** |␣
↪`{decision_config(row['config'])}` | "
        f"**{row['recall_at_5']:.4f}** | {row['latency_ms']:.3f} ms | "
        f"{row['speedup_vs_exact']:.2f}x | {row['candidate_fraction'] * 100:.
↪2f}% ({effort_label}) |"
    )

```

```

md_content = f"""
### Task 6: Outcome & Observations

```

The two required plots show the quality-efficiency trade-off across all eight␣
↪ANN configurations. The exact search is shown as the full-scan reference at␣
↪0% search efficiency and as the Recall@5 ceiling. The efficiency figure uses␣
↪a full-range panel plus an ANN zoom so both references and clustered ANN␣
↪points remain readable.

```

| Decision | Method | Configuration | Recall@5 | Average latency | Speed␣
↪relative to exact | Candidate fraction |
|---|---|---|---:|---:|---:|
{chr(10).join(decision_rows)}

```

****Best-balanced rule:**** Select the fastest ANN configuration whose recall is at␣
↪least 95% of the exact baseline. The resulting retention floor is␣
↪****{quality_floor:.4f}****, and the current selection is␣
↪****{balanced_selection_note.lower()}****

****Large-query deployment choice:**** {large_query_recommendation} If maximum␣
↪recall is more important than latency, select ****{highest_recall['method']}␣
↪with `{decision\_config(highest\_recall['config'])}`****, which achieves␣
↪Recall@5 = ****{highest_recall['avg_recall_at_5']:.4f}**** at␣
↪****{highest_recall['avg_latency_s'] * 1000:.3f} ms/query****.

****Interpretation:**** The fastest configuration minimizes query latency, the␣
↪highest-recall configuration approaches the exact quality ceiling, and the␣
↪balanced configuration applies an explicit quality requirement. HNSW is the␣
↪strongest family in this run because it provides the highest measured recall␣
↪with lower latency than IVF at comparable recall. IVF remains useful when␣
↪partition-based indexing is preferred, but its high-`nprobe` configurations␣
↪can be slower than exact search because of index overhead.

```

**Efficiency interpretation:** IVF candidate fractions use measured selected_
↳inverted-list sizes. HNSW candidate fractions use the `efSearch` search-work_
↳proxy because visited-node counts are not exposed by this FAISS interface.
↳These two effort values should therefore be interpreted as indicative, not_
↳directly equivalent.

**Generated plots:** `{os.path.basename(latency_plot_path)}` and `{os.path.
↳basename(eficiency_plot_path)}` in `{FIGURE_DIR}`.
"""

display(Markdown(md_content))

```

2.4.3 Task 6: Outcome & Observations

The two required plots show the quality-efficiency trade-off across all eight ANN configurations. The exact search is shown as the full-scan reference at 0% search efficiency and as the Recall@5 ceiling. The efficiency figure uses a full-range panel plus an ANN zoom so both references and clustered ANN points remain readable.

Decision	Method	Configuration	Recall@5	Average latency	Speed relative to exact	Candidate fraction
Global fastest	HNSW	efSearch=8	0.7020	0.031 ms	2.21x	0.08% (HNSW proxy)
Global highest recall	HNSW	efSearch=64	0.7320	0.090 ms	0.76x	0.64% (HNSW proxy)
Best bal- anced	HNSW	efSearch=8	0.7020	0.031 ms	2.21x	0.08% (HNSW proxy)

Best-balanced rule: Select the fastest ANN configuration whose recall is at least 95% of the exact baseline. The resulting retention floor is **0.6973**, and the current selection is **fastest configuration meeting the 95% exact-recall retention floor**.

Large-query deployment choice: For serving a large number of queries, select **HNSW with efSearch=8** as the default balanced configuration. It achieves Recall@5 = **0.7020** (95.6% of exact recall), average latency = **0.031 ms/query**, and **2.21x** speedup versus exact search. If maximum recall is more important than latency, select **HNSW with efSearch=64**, which achieves Recall@5 = **0.7320** at **0.090 ms/query**.

Interpretation: The fastest configuration minimizes query latency, the highest-recall configuration approaches the exact quality ceiling, and the balanced configuration applies an explicit quality requirement. HNSW is the strongest family in this run because it provides the highest measured recall with lower latency than IVF at comparable recall. IVF remains useful when partition-based indexing is preferred, but its high-*nprobe* configurations can be slower than exact search because of index overhead.

Efficiency interpretation: IVF candidate fractions use measured selected inverted-list sizes. HNSW candidate fractions use the `efSearch` search-work proxy because visited-node counts are not exposed by this FAISS interface. These two effort values should therefore be interpreted as indicative, not directly equivalent.

Generated plots: `task6_recall_vs_query_latency.png` and `task6_recall_vs_search_efficiency.png` in `data/figures_sciq`.

2.5 Module 4

2.6 Task 7: Qualitative Retrieval Analysis (2 Marks)

Compare Top-5 retrieval from **Embedding Model 1: all-MiniLM-L6-v2** and **Embedding Model 2: multi-qa-mpnet-base-dot-v1** on the same SciQ corpus and queries. Ten queries are selected deterministically to cover different query lengths and semantic regions, avoiding cherry-picking. The analysis reports Top-5 overlap, support-passage rank, and lexical overlap, allowing differences to be discussed in terms of **semantic similarity/ranking, keyword overlap, and short/long queries**.

```
[22]: # Task 7 - 10-query qualitative comparison

import numpy as np
import pandas as pd
import re
import os
import zipfile
from IPython.display import display

M1 = "all-MiniLM-L6-v2"
M2 = "multi-qa-mpnet-base-dot-v1"
K = 5

# =====
# Robustness: reload Task 1-2 artifacts if variables are absent
# =====

DATA_DIR = globals().get("DATA_DIR", "data/raw_sciq")
RESULTS_DIR = globals().get("RESULTS_DIR", "data/results_sciq")

os.makedirs(DATA_DIR, exist_ok=True)
os.makedirs(RESULTS_DIR, exist_ok=True)

# If required parquet files are not present, extract them from data.zip
required_files = [
    "corpus_sciq.parquet",
    "queries_sciq.parquet",
    "qrels_sciq.parquet"
]
```

```

if not all(
    os.path.exists(os.path.join(DATA_DIR, f))
    for f in required_files
):
    for zpath in ["data.zip", "/mnt/data/data.zip"]:

        if os.path.exists(zpath):

            with zipfile.ZipFile(zpath) as z:

                names = z.namelist()

                for target in required_files:

                    hits = [
                        n for n in names
                        if n.endswith(target)
                    ]

                    if hits:

                        src = hits[0]
                        out = os.path.join(DATA_DIR, target)

                        with z.open(src) as fi, open(out, "wb") as fo:
                            fo.write(fi.read())

                break

# =====
# Load datasets
# =====

if "corpus_df" not in globals():
    corpus_df = pd.read_parquet(
        os.path.join(DATA_DIR, "corpus_sciq.parquet")
    )

if "queries_df" not in globals():
    queries_df = pd.read_parquet(
        os.path.join(DATA_DIR, "queries_sciq.parquet")
    )

if "qrels_df" not in globals():
    qrels_df = pd.read_parquet(
        os.path.join(DATA_DIR, "qrels_sciq.parquet")
    )

```

```

)

# =====
# Load cached embeddings
# =====

if "embeddings_dict" not in globals():

    embeddings_dict = {}

    for m in [M1, M2]:

        p = os.path.join(
            RESULTS_DIR,
            f"emb_corpus_{m.replace('/', '__')}.npz"
        )

        if os.path.exists(p):
            embeddings_dict[m] = np.load(p)

if "query_embs_dict" not in globals():

    query_embs_dict = {}

    for m in [M1, M2]:

        p = os.path.join(
            RESULTS_DIR,
            f"emb_queries_{m.replace('/', '__')}.npz"
        )

        if os.path.exists(p):
            query_embs_dict[m] = np.load(p)

# =====
# Validate required artifacts
# =====

assert M1 in embeddings_dict and M2 in embeddings_dict, (
    "Run Tasks 1-2 first so both corpus embedding files exist."
)

assert M1 in query_embs_dict and M2 in query_embs_dict, (
    "Run Tasks 1-2 first so both query embedding files exist."
)

```

```

)

# =====
# Helper functions
# =====

def toks(s):
    return set(
        re.findall(
            r"[a-z0-9]+",
            str(s).lower()
        )
    )

def lex(q, p):

    a = toks(q)
    b = toks(p)

    return len(a & b) / len(a) if a else 0.0

def topk(model, qv):

    D = embeddings_dict[model].astype("float32")

    D /= np.linalg.norm(
        D,
        axis=1,
        keepdims=True
    ).clip(min=1e-12)

    q = qv.astype("float32")

    q /= np.linalg.norm(
        q,
        keepdims=True
    ).clip(min=1e-12)

    s = D @ q

    ix = np.argpartition(
        s,
        -K
    )[-K:]

```

```

    ix = ix[
        np.argsort(-s[ix])
    ]

    return ix, s[ix]

# =====
# Relevance mapping
# =====

qrels_map = (
    qrels_df
    .groupby("query_id")["doc_id"]
    .apply(set)
    .to_dict()
)

# =====
# Corpus lookup
# =====

corpus_lookup = (
    corpus_df
    .set_index("doc_id")["passage"]
    .to_dict()
)

doc_ids = corpus_df["doc_id"].tolist()

# =====
# Query information
# =====

texts = queries_df["text"].astype(str).tolist()

lengths = np.array([
    len(str(x).split())
    for x in texts
])

# =====
# Use Model 1 query embeddings to select representative queries

```

```

# =====

Q = query_embs_dict[M1].astype("float32")

Q /= np.linalg.norm(
    Q,
    axis=1,
    keepdims=True
).clip(min=1e-12)

# =====
# Select 10 representative queries
# =====

anchors = []

for q in np.linspace(.05, .95, 10):

    i = int(
        np.argmin(
            abs(
                lengths -
                np.quantile(lengths, q)
            )
        )
    )

    if i not in anchors:
        anchors.append(i)

selected = [anchors[0]]

while len(selected) < 10:

    rem = [
        i
        for i in range(len(texts))
        if i not in selected
    ]

    d = 1 - (
        Q[rem] @ Q[selected].T
    ).max(axis=1)

    target = np.quantile(

```

```

        lengths,
        len(selected) / 9
    )

    bonus = (
        .10 *
        (
            1 -
            np.minimum(
                abs(lengths[rem] - target) /
                max(lengths.max(), 1),
                1
            )
        )
    )

    selected.append(
        rem[
            int(np.argmax(d + bonus))
        ]
    )

selected = sorted(selected)

# =====
# Run Model 1 and Model 2 Top-5 retrieval
# =====

rows = []
details = []

summaries = {
    M1: [],
    M2: []
}

for i in selected:

    qid = queries_df.iloc[i]["query_id"]
    q = texts[i]

    rel = qrels_map.get(
        qid,
        set()
    )

```

```

)

model_ids = {}

# -----
# Retrieve using both embedding models
# -----

for m in [M1, M2]:

    ix, sc = topk(
        m,
        query_embs_dict[m][i]
    )

    ids = [
        doc_ids[int(j)]
        for j in ix
    ]

    model_ids[m] = ids

    # Find relevant document rank
    rank = next(
        (
            r
            for r, did in enumerate(ids, 1)
            if did in rel
        ),
        None
    )

    # Calculate lexical overlap for Top-5
    lx = [
        lex(
            q,
            corpus_lookup[d]
        )
        for d in ids
    ]

    summaries[m].append(
        (

```

```

        rank,
        lx[0]
    )
)

# Store individual Top-5 results
for r, (did, s, l) in enumerate(
    zip(ids, sc, lx),
    1
):

    passage_preview = (
        corpus_lookup[did]
        [:240]
        .replace("\n", " ")
        + "... "
    )

    details.append(
        [
            qid,
            m,
            r,
            did,
            round(float(s), 4),
            round(float(l), 3),
            passage_preview
        ]
    )

# -----
# Compare Model 1 vs Model 2
# -----

top5_overlap = len(
    set(model_ids[M1]) &
    set(model_ids[M2])
)

m1_support_rank = next(
    (
        r
        for r, d in enumerate(
            model_ids[M1],

```

```

        1
    )
    if d in rel
),
    "Not in Top-5"
)

m2_support_rank = next(
    (
        r
        for r, d in enumerate(
            model_ids[M2],
            1
        )
        if d in rel
    ),
    "Not in Top-5"
)

m1_lexical = round(
    summaries[M1][-1][1],
    3
)

m2_lexical = round(
    summaries[M2][-1][1],
    3
)

rows.append(
    [
        qid,
        q,
        int(lengths[i]),
        top5_overlap,
        m1_support_rank,
        m2_support_rank,
        m1_lexical,
        m2_lexical
    ]
)

```

```
# =====
```

```

# Create qualitative comparison table
# =====

qualitative_df = pd.DataFrame(
    rows,
    columns=[
        "Query ID",
        "Query",
        "Words",
        "Top-5 overlap",
        "M1 support rank",
        "M2 support rank",
        "M1 Top-1 lexical overlap",
        "M2 Top-1 lexical overlap"
    ]
)

# =====
# Categorize retrieval differences
# =====

def cat(r):

    # Significant surface-vocabulary (lexical) difference between the two models
    if abs(
        r["M1 Top-1 lexical overlap"] -
        r["M2 Top-1 lexical overlap"]
    ) >= 0.20:
        return "Keyword overlap"

    # Query length is the salient factor (short or long query)
    if (
        r["Words"] <= 7 or
        r["Words"] >= 18
    ):
        return "Short / long query"

    # Otherwise the models mainly differ in which passages they rank
    return "Semantic similarity / ranking"

qualitative_df["Difference type"] = (
    qualitative_df.apply(
        cat,
        axis=1
    )
)

```

```

)

# =====
# Detailed Top-5 table
# =====

top5_df = pd.DataFrame(
    details,
    columns=[
        "Query ID",
        "Model",
        "Rank",
        "Doc ID",
        "Cosine score",
        "Lexical overlap",
        "Passage preview"
    ]
)

# =====
# Display Task 7 results
# =====

print("=" * 80)
print("TASK 7 - QUALITATIVE RETRIEVAL ANALYSIS")
print("=" * 80)

print("\n--- 10 Representative Queries: Model 1 vs Model 2 ---")

display(
    qualitative_df
)

print("\n--- Top-5 Results for Each Model ---")

display(
    top5_df
)

# =====
# Summary statistics
#
# IMPORTANT:

```

```

# Use s["rank"] rather than s.rank.
# s.rank is interpreted by pandas as the DataFrame rank()
# method, which caused the previous AttributeError.
# =====

print("\n--- Model-level Qualitative Summary ---")

for m in [M1, M2]:

    s = pd.DataFrame(
        summaries[m],
        columns=[
            "rank",
            "lex"
        ]
    )

    support_rate = (
        s["rank"]
        .notna()
        .mean()
    )

    mean_support_rank = (
        s["rank"]
        .dropna()
        .mean()
    )

    mean_lexical_overlap = (
        s["lex"]
        .mean()
    )

    # Handle case where no relevant result occurs
    # in Top-5 for any query.
    if pd.isna(mean_support_rank):
        mean_rank_text = "N/A"
    else:
        mean_rank_text = f"{mean_support_rank:.2f}"

    print(
        f"{m}: "
        f"support in Top-5 = {support_rate:.1%}; "
        f"mean support rank = {mean_rank_text}; "

```

```

        f"mean Top-1 lexical overlap = "
        f"{mean_lexical_overlap:.3f}"
    )

# =====
# Difference-type summary
# =====

print("\n--- Retrieval Difference Types ---")

display(
    qualitative_df[
        "Difference type"
    ]
    .value_counts()
    .rename_axis("Difference type")
    .reset_index(name="Number of queries")
)

# =====
# Save Task 7 results
# =====

qualitative_df.to_csv(
    os.path.join(
        RESULTS_DIR,
        "task7_qualitative_query_comparison.csv"
    ),
    index=False
)

top5_df.to_csv(
    os.path.join(
        RESULTS_DIR,
        "task7_top5_results.csv"
    ),
    index=False
)

print("\nTask 7 completed successfully.")

print(
    "\nSaved:"
    "\n- task7_qualitative_query_comparison.csv"
)

```

```
"\n- task7_top5_results.csv"
)
```

```
=====
TASK 7 - QUALITATIVE RETRIEVAL ANALYSIS
=====
```

```
--- 10 Representative Queries: Model 1 vs Model 2 ---
```

Query ID	Query	Words	\
0 q_0003	What is the least dangerous radioactive decay?	7	
1 q_0026	A small scale version of what type of map disp...	13	
2 q_0029	Where do angiosperms produce seeds in flowers?	7	
3 q_0102	How many different amino acids make up proteins.	8	
4 q_0180	The atomic number of tungsten is 74. therefore...	37	
5 q_0276	Tests for levels of what in blood allow a diab...	18	
6 q_0316	Friction causes the molecules on rubbing surfa...	13	
7 q_0324	Electrical potential energy can be described b...	22	
8 q_0398	What term describes the orientation of a body ...	10	
9 q_0463	Light waves which interfere with themselves af...	18	

	Top-5 overlap	M1 support	rank	M2 support	rank	M1 Top-1 lexical overlap	\
0	2		1		1	0.857	
1	4		1		1	0.667	
2	4		2		1	0.429	
3	3	Not in Top-5			4	0.625	
4	2		1		1	0.895	
5	2		1	Not in Top-5		0.706	
6	4	Not in Top-5			3	0.154	
7	2		1		1	0.947	
8	1		1		1	0.818	
9	1		1		1	0.941	

	M2 Top-1 lexical overlap	Difference type
0	0.857	Short / long query
1	0.667	Semantic similarity / ranking
2	0.714	Keyword overlap
3	0.375	Keyword overlap
4	0.895	Short / long query
5	0.353	Keyword overlap
6	0.154	Semantic similarity / ranking
7	0.947	Short / long query
8	0.818	Semantic similarity / ranking
9	0.941	Short / long query

```
--- Top-5 Results for Each Model ---
```

Query ID	Model	Rank	Doc ID	Cosine score	\
----------	-------	------	--------	--------------	---

0	q_0003	all-MiniLM-L6-v2	1	doc_00003	0.7910
1	q_0003	all-MiniLM-L6-v2	2	doc_06010	0.6577
2	q_0003	all-MiniLM-L6-v2	3	doc_01506	0.6237
3	q_0003	all-MiniLM-L6-v2	4	doc_09445	0.6227
4	q_0003	all-MiniLM-L6-v2	5	doc_04967	0.6093
..	...	...	...	...	...
95	q_0463	multi-qa-mpnet-base-dot-v1	1	doc_00463	0.7627
96	q_0463	multi-qa-mpnet-base-dot-v1	2	doc_00182	0.6761
97	q_0463	multi-qa-mpnet-base-dot-v1	3	doc_05929	0.6603
98	q_0463	multi-qa-mpnet-base-dot-v1	4	doc_01241	0.6512
99	q_0463	multi-qa-mpnet-base-dot-v1	5	doc_03801	0.6486

	Lexical overlap	Passage preview
0	0.857	All radioactive decay is dangerous to living t...
1	0.857	All types of radioactive decay pose risks to l...
2	0.714	Radioactive decay is the process in which unst...
3	0.571	Radioactive decay is the breakdown of unstable...
4	0.571	When an atomic nucleus decays, it does so by r...
..	...	...
95	0.941	Waves are characterized by their ability to co...
96	0.118	Transmitted light may be refracted or scattere...
97	0.176	Light may be refracted by matter. The light is...
98	0.118	Diffraction occurs when waves spread out as th...
99	0.176	Three ways that waves may interact with matter...

[100 rows x 7 columns]

--- Model-level Qualitative Summary ---

all-MiniLM-L6-v2: support in Top-5 = 80.0%; mean support rank = 1.12; mean Top-1 lexical overlap = 0.704

multi-qa-mpnet-base-dot-v1: support in Top-5 = 90.0%; mean support rank = 1.56; mean Top-1 lexical overlap = 0.672

--- Retrieval Difference Types ---

	Difference type	Number of queries
0	Short / long query	4
1	Semantic similarity / ranking	3
2	Keyword overlap	3

Task 7 completed successfully.

Saved:

- task7_qualitative_query_comparison.csv
- task7_top5_results.csv

2.6.1 Task 7: Retrieval-Difference Analysis

Semantic similarity/ranking: Low Top-5 overlap means the models place different passages near the same query. A result can be semantically appropriate even when its wording differs from the query.

Keyword overlap: Query-to-passage lexical overlap shows how much surface vocabulary is shared. Higher overlap can help terminology-heavy questions, but lexical overlap alone does not establish semantic relevance.

Short/long queries: Short questions provide fewer contextual cues, while long questions contain more constraints and relationships. Different encoder representations can therefore produce different rankings.

multi-qa-mpnet-base-dot-v1 is the larger, QA-specialized 768-dimensional model, whereas all-MiniLM-L6-v2 is the smaller 384-dimensional model. The observed differences are interpreted as differences in the learned semantic spaces rather than simply differences in keyword matching.

2.7 Task 8: Final Recommendation (1 Mark)

The following recommendation is generated from the measured results in Tasks 2–7. It keeps the existing Task 1–6 implementation unchanged and uses the same L2-normalized retrieval space.

```
[23]: # Task 8 - Evidence-based final recommendation
import numpy as np, pandas as pd, os
from IPython.display import display, Markdown

M1 = "all-MiniLM-L6-v2"
M2 = "multi-qa-mpnet-base-dot-v1"

# Full-query exact Recall@5 for both models on the same L2-normalized space.
qmap = qrels_df.groupby("query_id")["doc_id"].apply(set).to_dict()
ids = corpus_df["doc_id"].tolist()
model_recall = {}
for m in [M1, M2]:
    D = embeddings_dict[m].astype("float32"); Q = query_embs_dict[m].
    ↪astype("float32")
    D /= np.linalg.norm(D, axis=1, keepdims=True).clip(min=1e-12)
    Q /= np.linalg.norm(Q, axis=1, keepdims=True).clip(min=1e-12)
    vals = []
    for i, r in queries_df.iterrows():
        s = D @ Q[i]; ix = np.argsort(s, -5)[-5:]; ix = ix[np.
        ↪argsort(-s[ix])]
        got = {ids[int(j)] for j in ix}; rel = qmap.get(r["query_id"], set())
        vals.append(len(got & rel) / len(rel) if rel else 0)
    model_recall[m] = float(np.mean(vals))

quality_model = M2 if model_recall[M2] >= model_recall[M1] else M1
resource_model = M1 if model_recall[M1] >= 0.95 * model_recall[M2] else M2
other = M1 if quality_model == M2 else M2
```

```

# ANN configurations chosen in Task 6 (dynamic - never hardcoded, so the text
# always matches the actual results when the notebook is re-run on any machine).
bcfg = balanced["config"]; bmethod = balanced["method"]
brec = float(balanced["avg_recall_at_5"]); blat =
    float(balanced["avg_latency_s"] * 1000)
bsp = float(balanced["speedup_vs_exact"]); retention = brec /
    exact_baseline_recall
fcfg = fastest["config"]; frec = float(fastest["avg_recall_at_5"]); flat =
    float(fastest["avg_latency_s"] * 1000)
hcfg = highest_recall["config"]
enc = {r["Model Name"]: r["Runtime (s)"] for r in model_metrics}
enc_m1 = enc.get(M1, float("nan")); enc_m2 = enc.get(M2, float("nan"))

recommendation_table = pd.DataFrame([
    ["Embedding model", quality_model, round(model_recall[quality_model], 4),
    "Higher full-query Recall@5"],
    ["Similarity metric", "Cosine / Inner Product after L2 normalization",
    "Rank-equivalent", "Task 3: 0/25 normalized ranking differences"],
    ["Best ANN trade-off", f"{bmethod} {bcfg}", round(brec, 4), f"{blat:.3f} ms/
    query; {bsp:.2f}x exact speed"],
    ["Resource-constrained", f"{resource_model} + {bmethod} {fcfg}",
    round(model_recall[resource_model], 4), "384-D MiniLM and fastest ANN
    setting"],
], columns=["Decision", "Selection", "Measured evidence", "Reason"])
display(recommendation_table)

md = f"""### Task 8: Final Recommendation

**1. Better embedding model - {quality_model}.** Full-query exact Recall@5 is
    {model_recall[quality_model]:.4f}, compared with {model_recall[other]:.
    4f} for {other}. Task 7 provides qualitative Top-5/ranking evidence for 10
    representative queries.

**2. Most appropriate similarity metric - Cosine similarity, implemented as
    inner product after L2 normalization.** Task 3 verified identical rankings
    for normalized Cosine, Dot Product and L2.

**3. Best ANN trade-off - {bmethod} {bcfg}.** Task 6 selected this as the
    fastest configuration satisfying the 95%-of-exact-recall rule. It gives
    Recall@5 {brec:.4f}, about {retention*100:.1f}% of exact Recall@5
    {exact_baseline_recall:.4f}, at {blat:.3f} ms/query and {bsp:.
    2f}x exact-search speed.

```

```

**4. Resource-constrained deployment - {resource_model} + {bmethod} `{fcfg}`.**
↳ {M1} uses 384 dimensions versus {M2}'s 768, and Task 2 measured {enc_m1:.2f}
↳ s versus {enc_m2:.2f} s to encode the corpus and query texts. `{fcfg}` is
↳ the fastest ANN configuration: **{flat:.3f} ms/query** at Recall@5 = **{frec:
↳ .4f}**. If this quality is insufficient, move to the balanced setting
↳ `{bcfg}`.

**5. Main limitation -** the experiment uses 10,000 passages and 500 SciQ
↳ queries, so results may not transfer directly to much larger or
↳ domain-shifted production data. HNSW candidate effort is also an `efSearch`
↳ proxy rather than a directly exposed visited-node count in this FAISS setup.

**6. Production deployment -** use **offline document ingestion ->
↳ {resource_model} embedding -> L2 normalization -> versioned HNSW index ->
↳ query API -> query embedding + normalization -> ANN Top-K retrieval ->
↳ optional re-ranking/response generation**. Keep document embedding offline;
↳ only query embedding runs online. Monitor p50/p95/p99 latency, Recall@5 on a
↳ labelled regression set, index size, failures and drift. Rebuild/version the
↳ index when the corpus or model changes, and rerun the Task 7 qualitative
↳ suite before release.

### Overall decision
**{quality_model} for retrieval quality (highest full-query Recall@5),
↳ {resource_model} for constrained resources, and {bmethod} `{bcfg}` as the
↳ default balanced ANN configuration.** Use `{fcfg}` for the lowest-resource
↳ case and `{hcfg}` when maximum measured recall is the priority. Modules 2-3
↳ use {M2} as the single selected model for the similarity and ANN study,
↳ while this recommendation uses the full two-model Recall@5 comparison.
"""
display(Markdown(md))
recommendation_table.to_csv(os.path.join(RESULTS_DIR,
↳ "task8_final_recommendation.csv"), index=False)

```

	Decision	Selection \
0	Embedding model	all-MiniLM-L6-v2
1	Similarity metric	Cosine / Inner Product after L2 normalization
2	Best ANN trade-off	HNSW {'efSearch': 8}
3	Resource-constrained	all-MiniLM-L6-v2 + HNSW {'efSearch': 8}

	Measured evidence	Reason
0	0.744	Higher full-query Recall@5
1	Rank-equivalent	Task 3: 0/25 normalized ranking differences
2	0.702	0.031 ms/query; 2.21x exact speed
3	0.744	384-D MiniLM and fastest ANN setting

2.7.1 Task 8: Final Recommendation

1. Better embedding model - all-MiniLM-L6-v2. Full-query exact Recall@5 is **0.7440**, compared with **0.7340** for multi-qa-mpnet-base-dot-v1. Task 7 provides qualitative Top-5/ranking evidence for 10 representative queries.

2. Most appropriate similarity metric - Cosine similarity, implemented as inner product after L2 normalization. Task 3 verified identical rankings for normalized Cosine, Dot Product and L2.

3. Best ANN trade-off - HNSW {'efSearch': 8}. Task 6 selected this as the fastest configuration satisfying the 95%-of-exact-recall rule. It gives Recall@5 **0.7020**, about **95.6%** of exact Recall@5 **0.7340**, at **0.031 ms/query** and **2.21x** exact-search speed.

4. Resource-constrained deployment - all-MiniLM-L6-v2 + HNSW {'efSearch': 8}. all-MiniLM-L6-v2 uses 384 dimensions versus multi-qa-mpnet-base-dot-v1's 768, and Task 2 measured 12.44 s versus 87.59 s to encode the corpus and query texts. {'efSearch': 8} is the fastest ANN configuration: **0.031 ms/query** at Recall@5 = **0.7020**. If this quality is insufficient, move to the balanced setting {'efSearch': 8}.

5. Main limitation - the experiment uses 10,000 passages and 500 SciQ queries, so results may not transfer directly to much larger or domain-shifted production data. HNSW candidate effort is also an efSearch proxy rather than a directly exposed visited-node count in this FAISS setup.

6. Production deployment - use offline document ingestion -> all-MiniLM-L6-v2 embedding -> L2 normalization -> versioned HNSW index -> query API -> query embedding + normalization -> ANN Top-K retrieval -> optional re-ranking/response generation. Keep document embedding offline; only query embedding runs online. Monitor p50/p95/p99 latency, Recall@5 on a labelled regression set, index size, failures and drift. Rebuild/version the index when the corpus or model changes, and rerun the Task 7 qualitative suite before release.

2.7.2 Overall decision

all-MiniLM-L6-v2 for retrieval quality (highest full-query Recall@5), all-MiniLM-L6-v2 for constrained resources, and HNSW {'efSearch': 8} as the default balanced ANN configuration. Use {'efSearch': 8} for the lowest-resource case and {'efSearch': 64} when maximum measured recall is the priority. Modules 2-3 use multi-qa-mpnet-base-dot-v1 as the single selected model for the similarity and ANN study, while this recommendation uses the full two-model Recall@5 comparison.